

Analysing European Crop Production Trends with FAOSTAT Dataset

Brief Summary

This notebook explores agricultural trends in Europe by analyzing FAOSTAT data spanning 1961–2023 for 144 major crops. Through an iterative “Pytudes” -style workflow, we will:

1. **Load & Clean** the raw CSV data.
2. **Transform** it into time-series and normalized metrics (e.g., yield per hectare).
3. **Visualize** aggregate and crop-specific trends, highlighting key shifts and anomalies.
4. **Compare** subregions (Western, Eastern, Northern, Southern Europe) to uncover differential growth patterns.
5. **Experiment** with advanced chart types (heatmaps, small multiples, animated race charts) to deepen insights.

Our goal is to demonstrate a reproducible, critique-driven data-visualization process that reveals how European crop production, yields, and harvested areas have evolved over six decades.

Here lie all the imports needed across this exploration

```
In [88]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import textwrap
import squarify
import matplotlib.ticker as ticker
from matplotlib.colors import LogNorm
```

So as we needed to download the data in 3 separate files (due to site limitations) we want to have all the variable in one place for more easier use, so we are gonna load the 3 datasets, merge them and then clean some of the null values.

```
In [62]: # 1. Load each CSV separately
df_area = pd.read_csv('FAOSTAT_data_en_5-31-2025-harvest-area.csv')
df_prod = pd.read_csv('FAOSTAT_data_en_5-31-2025-production.csv')
```

```
df_yld = pd.read_csv('FAOSTAT_data_en_5-31-2025-yield.csv')

# 2. Inspect columns to identify the common keys (including 'Unit')
print("Area-harvest columns:", df_area.columns.tolist())
print("Production    columns:", df_prod.columns.tolist())
print("Yield          columns:", df_yld.columns.tolist())

# 3. Reduce each DataFrame to the key columns + its own 'Value' and 'Unit', renaming as needed.
#   Keep: ['Area', 'Item', 'Year', 'Value', 'Unit'] → rename Value and Unit
df_area_trim = (
    df_area
    .loc[:, ['Area', 'Item', 'Year', 'Value', 'Unit']]
    .rename(columns={
        'Value': 'Area_harvested',
        'Unit': 'Unit_Area_harvested'
    })
)

df_prod_trim = (
    df_prod
    .loc[:, ['Area', 'Item', 'Year', 'Value', 'Unit']]
    .rename(columns={
        'Value': 'Production',
        'Unit': 'Unit_Production'
    })
)

df_yld_trim = (
    df_yld
    .loc[:, ['Area', 'Item', 'Year', 'Value', 'Unit']]
    .rename(columns={
        'Value': 'Yield',
        'Unit': 'Unit_Yield'
    })
)

# 4. Merge the three DataFrames on the common keys: ['Area', 'Item', 'Year']
#   Use outer merge so no data is lost if one file is missing a combination.
df_merge1 = pd.merge(
    df_area_trim,
    df_prod_trim,
    on=['Area', 'Item', 'Year'],
```

```
        how='outer'
    )

df_merged = pd.merge(
    df_merge1,
    df_yld_trim,
    on=['Area', 'Item', 'Year'],
    how='outer'
)

# 5. Check the result structure
print("Merged shape:", df_merged.shape)
print(df_merged.sample(10))

# 6. (Optional) Drop rows where all three metrics are NaN
df_merged = df_merged.dropna(subset=['Area_harvested', 'Production', 'Yield'], how='all')

# 7. Inspect how many non-null entries exist, and confirm units
print("\nNon-null counts after merging:")
print(df_merged[['Area_harvested', 'Production', 'Yield']].notna().sum())

print("\nUnique units for each metric:")
print(" Area harvested units:", df_merged['Unit_Area_harvested'].unique())
print(" Production units:   ", df_merged['Unit_Production'].unique())
print(" Yield units:       ", df_merged['Unit_Yield'].unique())
```

```

Area-harvest columns: ['Domain', 'Area', 'Element', 'Item', 'Year', 'Unit', 'Value']
Production columns: ['Domain', 'Area', 'Element', 'Item', 'Year', 'Unit', 'Value']
Yield columns: ['Domain', 'Area', 'Element', 'Item', 'Year', 'Unit', 'Value']
Merged shape: (150724, 9)

```

	Area \
32573	Denmark
37070	Finland
123003	Spain
112817	Romania
98114	Poland
99432	Portugal
16183	Bosnia and Herzegovina
112425	Romania
24397	Croatia
2636	Albania

	Item Year \
32573	Linseed 1998
37070	Carrots and turnips 1999
123003	Anise, badian, coriander, cumin, caraway, fenn... 1992
112817	Watermelons 2006
98114	Strawberries 1998
99432	Broad beans and horse beans, dry 1979
16183	Lemons and limes 2013
112425	Sunflower-seed oil, crude 1997
24397	Eggplants (aubergines) 2016
2636	Other fruits, n.e.c. 2023

	Area harvested	Unit_Area harvested	Production	Unit_Production \
32573	3871.0	ha	1645.0	t
37070	1726.0	ha	62309.0	t
123003	2242.0	ha	1055.0	t
112817	30689.0	ha	587756.0	t
98114	52614.0	ha	149858.0	t
99432	37700.0	ha	21300.0	t
16183	1.0	ha	12.0	t
112425	NaN	NaN	375000.0	t
24397	NaN	ha	NaN	t
2636	150.0	ha	3840.8	t

	Yield	Unit_Yield
32573	425.0	kg/ha

```
37070    36100.2      kg/ha
123003     470.6      kg/ha
112817   19152.0      kg/ha
98114     2848.3      kg/ha
99432      565.0      kg/ha
16183    12000.0      kg/ha
112425       NaN       NaN
24397       NaN       NaN
2636     25656.6      kg/ha
```

Non-null counts after merging:

```
Area_harvested    103634
Production        134279
Yield             97177
dtype: int64
```

Unique units for each metric:

```
Area harvested units: ['ha' 'nan']
Production units:     ['t' 'nan']
Yield units:          [nan 'kg/ha']
```

```
In [63]: df_merged.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 134538 entries, 0 to 150723
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Area              134538 non-null   object 
 1   Item              134538 non-null   object 
 2   Year              134538 non-null   int64  
 3   Area_harvested    103634 non-null   float64
 4   Unit_Area_harvested 108660 non-null   object 
 5   Production        134279 non-null   float64
 6   Unit_Production    134346 non-null   object 
 7   Yield             97177 non-null    float64
 8   Unit_Yield         97177 non-null   object 
dtypes: float64(3), int64(1), object(5)
memory usage: 10.3+ MB
```

As this dataset has a big number of entries, we would like to see exactly how many types of items we have, basically seeing which crops or products we cover in this analysis.

```
In [64]: unique_items = df_merged['Item'].unique()
print(len(unique_items))
for item in sorted(unique_items):
    print(item)
```

Almonds, in shell
Anise, badian, coriander, cumin, caraway, fennel and juniper berries, raw
Apples
Apricots
Artichokes
Asparagus
Avocados
Bananas
Barley
Beans, dry
Beer of barley, malted
Blueberries
Broad beans and horse beans, dry
Broad beans and horse beans, green
Buckwheat
Cabbages
Canary seed
Cantaloupes and other melons
Carrots and turnips
Castor oil seeds
Cauliflowers and broccoli
Cereals n.e.c.
Cherries
Chestnuts, in shell
Chick peas, dry
Chicory roots
Chillies and peppers, dry (*Capsicum spp.*, *Pimenta spp.*), raw
Chillies and peppers, green (*Capsicum spp.* and *Pimenta spp.*)
Coconut oil
Coconuts, in shell
Coffee, green
Cotton lint, ginned
Cotton seed
Cottonseed oil
Cow peas, dry
Cranberries
Cucumbers and gherkins
Currants
Dates
Edible roots and tubers with high starch or inulin content, n.e.c., fresh
Eggplants (aubergines)

Figs
Flax, raw or retted
Gooseberries
Grapes
Green corn (maize)
Green garlic
Green tea (not fermented), black tea (fermented) and partly fermented tea, in immediate packings of a content not exceeding 3 kg
Groundnut oil
Groundnuts, excluding shelled
Hazelnuts, in shell
Hempseed
Hop cones
Kenaf, and other textile bast fibres, raw or retted
Kiwi fruit
Leeks and other alliaceous vegetables
Lemons and limes
Lentils, dry
Lettuce and chicory
Linseed
Locust beans (carobs)
Lupins
Maize (corn)
Margarine and shortening
Melonseed
Millet
Mixed grain
Molasses
Mushrooms and truffles
Mustard seed
Oats
Oil of linseed
Oil of maize
Oil of palm kernel
Oil of sesame seed
Okra
Olive oil
Olives
Onions and shallots, dry (excluding dehydrated)
Onions and shallots, green
Oranges
Other beans, green
Other berries and fruits of the genus *vaccinium* n.e.c.

Other citrus fruit, n.e.c.
Other fruits, n.e.c.
Other nuts (excluding wild edible nuts and groundnuts), in shell, n.e.c.
Other oil seeds, n.e.c.
Other pome fruits
Other pulses n.e.c.
Other stimulant, spice and aromatic crops, n.e.c.
Other stone fruits
Other tropical fruits, n.e.c.
Other vegetables, fresh n.e.c.
Palm oil
Peaches and nectarines
Pears
Peas, dry
Peas, green
Peppermint, spearmint
Persimmons
Pineapples
Pistachios, in shell
Plums and sloes
Pomelos and grapefruits
Poppy seed
Potatoes
Pumpkins, squash and gourds
Pyrethrum, dried flowers
Quinces
Rape or colza seed
Rapeseed or canola oil, crude
Raspberries
Raw cane or beet sugar (centrifugal only)
Rice
Rye
Safflower seed
Safflower-seed oil, crude
Seed cotton, unginned
Sesame seed
Sorghum
Sour cherries
Soya bean oil
Soya beans
Spinach
Strawberries

String beans
Sugar beet
Sugar cane
Sunflower seed
Sunflower-seed oil, crude
Sweet potatoes
Tangerines, mandarins, clementines
Tea leaves
Tomatoes
Triticale
True hemp, raw or retted
Tung nuts
Unmanufactured tobacco
Vetches
Walnuts, in shell
Watermelons
Wheat
Wine
Yams

Now having this list we can see that we have 144 individual items, plotting 144 crops on any of the properties we have production area harvested or yield will be confusing and hard to read. Now to use some generative AI magic we create a map of all the crops we have to a more general list of categories

```
In [65]: crop_to_group = {
    'Almonds, in shell': 'Nuts',
    'Anise, badian, coriander, cumin, caraway, fennel and juniper berries, raw': 'Spices & Aromatics',
    'Apples': 'Fruits',
    'Apricots': 'Fruits',
    'Artichokes': 'Vegetables',
    'Asparagus': 'Vegetables',
    'Avocados': 'Fruits',
    'Bananas': 'Fruits',
    'Barley': 'Cereals',
    'Beans, dry': 'Pulses/Legumes',
    'Beer of barley, malted': 'Beverages',
    'Blueberries': 'Fruits',
    'Broad beans and horse beans, dry': 'Pulses/Legumes',
    'Broad beans and horse beans, green': 'Vegetables',
    'Buckwheat': 'Cereals',
```

'Cabbages':	'Vegetables',
'Canary seed':	'Cereals',
'Cantaloupes and other melons':	'Fruits',
'Carrots and turnips':	'Roots & Tubers',
'Castor oil seeds':	'Oilseeds',
'Cauliflowers and broccoli':	'Vegetables',
'Cereals n.e.c.':	'Cereals',
'Cherries':	'Fruits',
'Chestnuts, in shell':	'Nuts',
'Chick peas, dry':	'Pulses/Legumes',
'Chicory roots':	'Vegetables',
'Chillies and peppers, dry (<i>Capsicum</i> spp., <i>Pimenta</i> spp.), raw':	'Spices & Aromatics',
'Chillies and peppers, green (<i>Capsicum</i> spp. and <i>Pimenta</i> spp.)':	'Vegetables',
'Coconut oil':	'Oils & Fats',
'Coconuts, in shell':	'Fruits',
'Coffee, green':	'Beverages',
'Cotton lint, ginned':	'Fiber Crops',
'Cotton seed':	'Oilseeds',
'Cottonseed oil':	'Oils & Fats',
'Cow peas, dry':	'Pulses/Legumes',
'Cranberries':	'Fruits',
'Cucumbers and gherkins':	'Vegetables',
'Currants':	'Fruits',
'Dates':	'Fruits',
'Edible roots and tubers with high starch or inulin content, n.e.c., fresh':	'Roots & Tubers',
'Eggplants (aubergines)':	'Vegetables',
'Figs':	'Fruits',
'Flax, raw or retted':	'Fiber Crops',
'Gooseberries':	'Fruits',
'Grapes':	'Fruits',
'Green corn (maize)':	'Cereals',
'Green garlic':	'Vegetables',
'Green tea (not fermented), black tea (fermented) and partly fermented tea, in immediate packings of a content not exceeding 3% moisture':	'Beverages',
'Groundnut oil':	'Oils & Fats',
'Groundnuts, excluding shelled':	'Oilseeds',
'Hazelnuts, in shell':	'Nuts',
'Hempseed':	'Oilseeds',
'Hop cones':	'Beverages',

'Kenaf, and other textile bast fibres, raw or retted':	
'Kiwi fruit':	'Fiber Crops',
'Leeks and other alliaceous vegetables':	'Fruits',
'Lemons and limes':	'Vegetables',
'Lentils, dry':	'Fruits',
'Lettuce and chicory':	'Pulses/Legumes',
'Linseed':	'Vegetables',
'Locust beans (carobs)':	'Oilseeds',
'Lupins':	'Vegetables',
'Maize (corn)':	'Pulses/Legumes',
'Margarine and shortening':	'Cereals',
'Melonseed':	'Oils & Fats',
'Millet':	'Oilseeds',
'Mixed grain':	'Cereals',
'Molasses':	'Sugar Crops',
'Mushrooms and truffles':	'Vegetables',
'Mustard seed':	'Oilseeds',
'Oats':	'Cereals',
'Oil of linseed':	'Oils & Fats',
'Oil of maize':	'Oils & Fats',
'Oil of palm kernel':	'Oils & Fats',
'Oil of sesame seed':	'Oils & Fats',
'Okra':	'Vegetables',
'Olive oil':	'Oils & Fats',
'Olives':	'Fruits',
'Onions and shallots, dry (excluding dehydrated)':	'Vegetables',
'Onions and shallots, green':	'Vegetables',
'Oranges':	'Fruits',
'Other beans, green':	'Vegetables',
'Other berries and fruits of the genus vaccinium n.e.c.':	
	'Fruits',
'Other citrus fruit, n.e.c.':	'Fruits',
'Other fruits, n.e.c.':	'Fruits',
'Other nuts (excluding wild edible nuts and groundnuts), in shell, n.e.c.':	
	'Nuts',
'Other oil seeds, n.e.c.':	'Oilseeds',
'Other pome fruits':	'Fruits',
'Other pulses n.e.c.':	'Pulses/Legumes',
'Other stimulant, spice and aromatic crops, n.e.c.':	
	'Spices & Aromatics',
'Other stone fruits':	'Fruits',

'Other tropical fruits, n.e.c.':	'Fruits',
'Other vegetables, fresh n.e.c.':	'Vegetables',
'Palm oil':	'Oils & Fats',
'Peaches and nectarines':	'Fruits',
'Pears':	'Fruits',
'Peas, dry':	'Pulses/Legumes',
'Peas, green':	'Vegetables',
'Peppermint, spearmint':	'Spices & Aromatics',
'Persimmons':	'Fruits',
'Pineapples':	'Fruits',
'Pistachios, in shell':	'Nuts',
'Plums and sloes':	'Fruits',
'Pomelos and grapefruits':	'Fruits',
'Poppy seed':	'Oilseeds',
'Potatoes':	'Roots & Tubers',
'Pumpkins, squash and gourds':	'Vegetables',
'Pyrethrum, dried flowers':	'Industrial Crops',
'Quinces':	'Fruits',
'Rape or colza seed':	'Oilseeds',
'Rapeseed or canola oil, crude':	'Oils & Fats',
'Raspberries':	'Fruits',
'Raw cane or beet sugar (centrifugal only)':	'Sugar Crops',
'Rice':	'Cereals',
'Rye':	'Cereals',
'Safflower seed':	'Oilseeds',
'Safflower-seed oil, crude':	'Oils & Fats',
'Seed cotton, unginned':	'Fiber Crops',
'Sesame seed':	'Oilseeds',
'Sorghum':	'Cereals',
'Sour cherries':	'Fruits',
'Soya bean oil':	'Oils & Fats',
'Soya beans':	'Pulses/Legumes',
'Spinach':	'Vegetables',
'Strawberries':	'Fruits',
'String beans':	'Vegetables',
'Sugar beet':	'Sugar Crops',
'Sugar cane':	'Sugar Crops',
'Sunflower seed':	'Oilseeds',
'Sunflower-seed oil, crude':	'Oils & Fats',
'Sweet potatoes':	'Roots & Tubers',
'Tangerines, mandarins, clementines':	'Fruits',
'Tea leaves':	'Beverages',

```

'Tomatoes':
'Triticale':
'True hemp, raw or retted':
'Tung nuts':
'Unmanufactured tobacco':
'Vetches':
'Walnuts, in shell':
'Watermelons':
'Wheat':
'Wine':
'Yams':
}

#To verify that every one of your 144 items is covered:
assert len(crop_to_group) == 144

```

Now lets try a simple visualization by analysing the production of the top 20 crops in our dataset

```

In [71]: # Because we have multiple years and countries and we are currently looking at total production of the crops we need to aggregate the
production_per_item = (
    df_merged
    .groupby('Item')['Production']
    .sum()
    .sort_values(ascending=False)
)

# 2. To compute total production by Group without altering df_merged, create a temporary DataFrame:
df_temp = df_merged.copy()
df_temp['Group'] = df_temp['Item'].map(crop_to_group)

production_by_group = (
    df_temp
    .groupby('Group')['Production']
    .sum()
    .sort_values(ascending=False)
)

# -----
# (A) Pie chart of all 144 crop items
# -----
fig1, ax1 = plt.subplots(figsize=(10, 10))

```

```

# Generate 144 distinct colors from a colormap
colors_items = plt.cm.tab20b(np.linspace(0, 1, len(production_per_item)))

wedges1, texts1, autotexts1 = ax1.pie(
    production_per_item.values,
    labels=production_per_item.index,
    autopct='%.1f%%',
    colors=colors_items,
    startangle=90,
    wedgeprops=dict(edgecolor='white', linewidth=0.5)
)

# Make labels very small because there are 144 slices
for txt in texts1:
    txt.set_fontsize(4)
for atxt in autotexts1:
    atxt.set_fontsize(3)

ax1.set_title('Total Production by Crop Item (all 144)', fontsize=14, pad=20)
plt.tight_layout()

# -----
# (B) Pie chart of grouped categories, with Legend (no on-slice labels)
# -----
fig2, ax2 = plt.subplots(figsize=(12, 12))

groups = production_by_group.index
values = production_by_group.values

# Compute percentages for each group
total = values.sum()
percentages = values / total * 100
legend_labels = [f'{grp}: {p:.3f}%' for grp, p in zip(groups, percentages)]

# Use a discrete colormap with at least Len(groups) distinct colors
cmap = plt.get_cmap('tab20')
colors_groups = cmap.colors[:len(groups)]

# Create the pie WITHOUT on-slice labels or percentages
wedges2, _ = ax2.pie(

```

```
values,
labels=None,           # No labels on the slices
startangle=90,
colors=colors_groups,
wedgeprops=dict(edgecolor='white', linewidth=1)
)

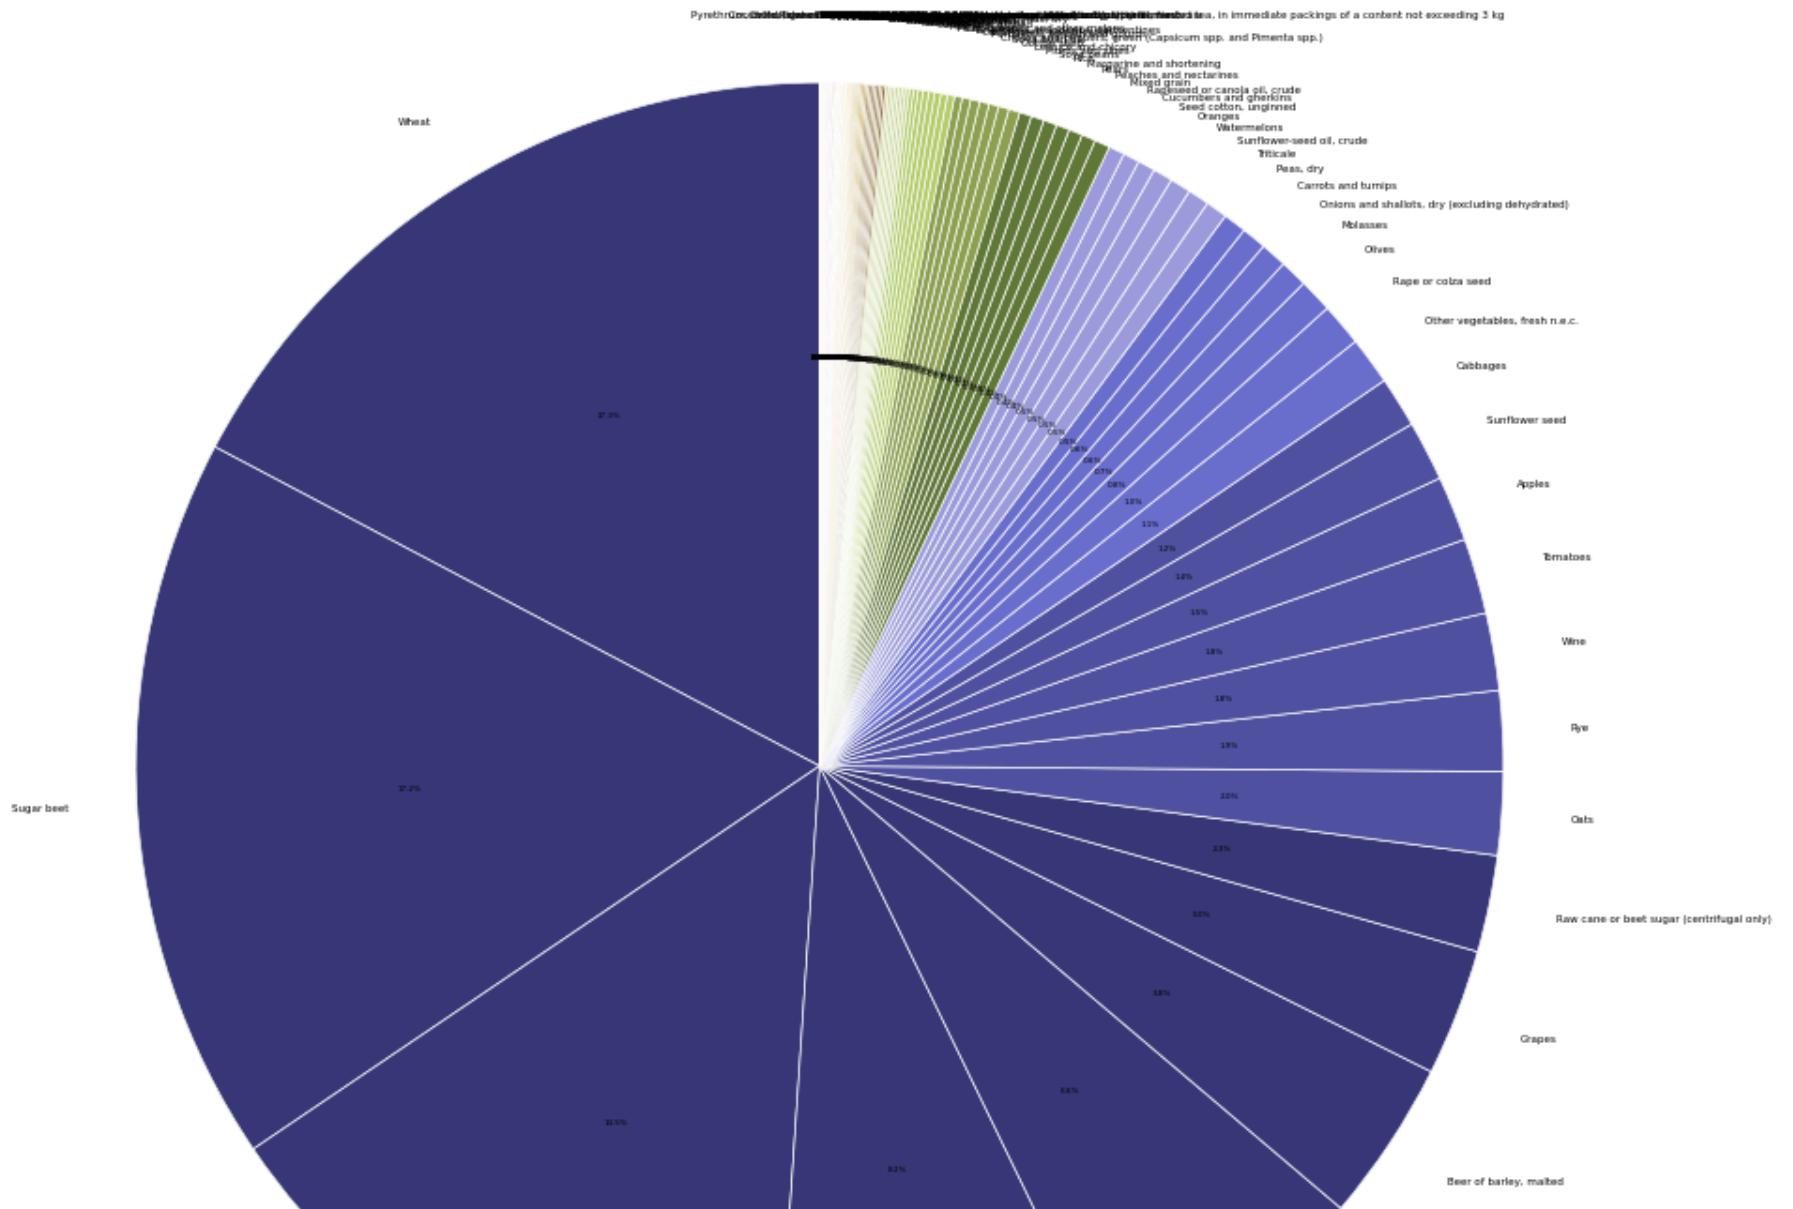
# Add a white circle for donut style (optional)
centre_circle = plt.Circle((0, 0), 0.45, color='white', linewidth=0)
ax2.add_artist(centre_circle)

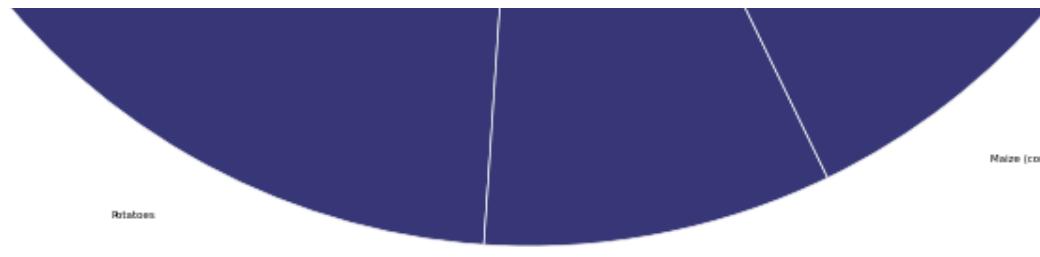
# Title
ax2.set_title('Total Production by Category (grouped)', fontsize=16, pad=20)

# Add Legend on the right with category names + percentages
ax2.legend(
    wedges2,
    legend_labels,
    title="Category - % of Total",
    loc="center left",
    bbox_to_anchor=(1, 0, 0.3, 1),   # Positions Legend to the right of the pie
    fontsize=10,
    title_fontsize=12
)

plt.tight_layout()
plt.show()
```

Total Production by Crop Item (all 144)



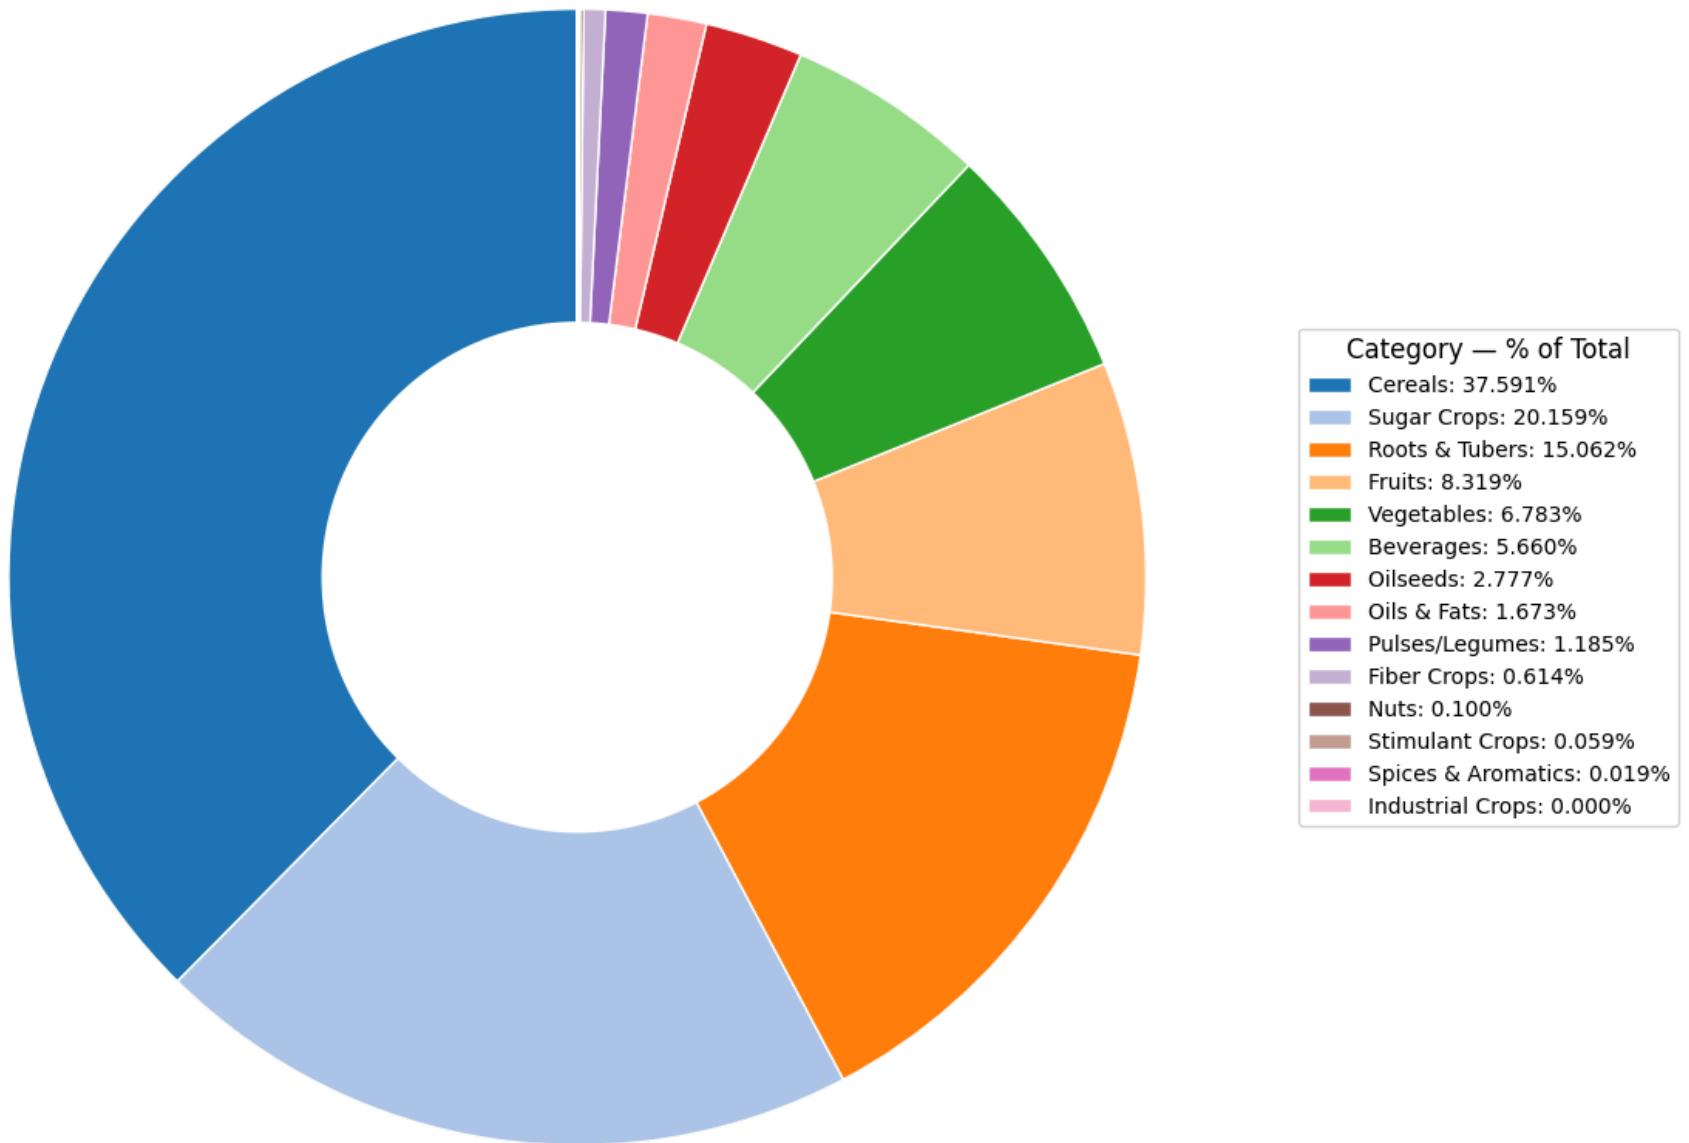


Barley

Potatoes

Maize (corn)

Total Production by Category (grouped)



As we can clearly see in the second visualization our pie chart already starts to look better, while even the grouped categories seem a bit cluttered at the lower end where the percentages are relatively small to our leader it can offer a better picture in how the dataset is spread across the types of crops that we have.

now lets do the same for area harvested and yield

```
In [72]: # -----
# (B) AREA_HARVESTED PIES (exactly as before)
# -----
```

```
# 1B. Aggregate total Area_Harvested by Item
area_per_item = (
    df_merged
    .groupby('Item')['Area_harvested']
    .sum()
    .sort_values(ascending=False)
)

# 2B. Aggregate total Area_Harvested by Group (using a temporary copy)
df_temp_area = df_merged.copy()
df_temp_area['Group'] = df_temp_area['Item'].map(crop_to_group)
area_by_group = (
    df_temp_area
    .groupby('Group')['Area_harvested']
    .sum()
    .sort_values(ascending=False)
)

# 3B. Plot: (B1) Pie of all 144 Items by Area_Harvested
figB1, axB1 = plt.subplots(figsize=(10, 10))
colors_items_B = plt.cm.tab20b(np.linspace(0, 1, len(area_per_item)))
wedgesB1, textsB1, autotextsB1 = axB1.pie(
    area_per_item.values,
    labels=area_per_item.index,
    autopct='%1.1f%',
    colors=colors_items_B,
```

```

        startangle=90,
        wedgeprops=dict(edgecolor='white', linewidth=0.5)
    )
    for txt in textsB1:
        txt.set_fontsize(4)
    for atxt in autotextsB1:
        atxt.set_fontsize(3)
    axB1.set_title('Total Area Harvested by Crop Item (all 144)', fontsize=14, pad=20)
    plt.tight_layout()

# 4B. Plot: (B2) Pie of Groups by Area_Harvested (legend on right)
figB2, axB2 = plt.subplots(figsize=(12, 8))
groups_B = area_by_group.index
values_B = area_by_group.values
total_B = values_B.sum()
pct_B = values_B / total_B * 100
legend_labels_B = [f"{grp}: {p:.3f}%" for grp, p in zip(groups_B, pct_B)]
cmap_B = plt.get_cmap('tab20')
colors_groups_B = cmap_B.colors[: len(groups_B)]
wedgesB2, _ = axB2.pie(
    values_B,
    labels=None,
    startangle=90,
    colors=colors_groups_B,
    wedgeprops=dict(edgecolor='white', linewidth=1)
)
centreB2 = plt.Circle((0, 0), 0.45, color='white', linewidth=0)
axB2.add_artist(centreB2)
axB2.set_title('Total Area Harvested by Category (grouped)', fontsize=16, pad=20)
axB2.legend(
    wedgesB2,
    legend_labels_B,
    title="Category - % of Total Area",
    loc="center left",
    bbox_to_anchor=(1, 0, 0.3, 1),
    fontsize=10,
    title_fontsize=12
)
plt.tight_layout()

# -----

```

```

# (C) YIELD PIES (using the same naming as AREA_HARVESTED)
# -----
# 1C. Aggregate total Yield by Item
yield_per_item = (
    df_merged
    .groupby('Item')['Yield']
    .sum()
    .sort_values(ascending=False)
)

# 2C. Aggregate total Yield by Group (using a temporary copy)
df_temp_yield = df_merged.copy()
df_temp_yield['Group'] = df_temp_yield['Item'].map(crop_to_group)
yield_by_group = (
    df_temp_yield
    .groupby('Group')['Yield']
    .sum()
    .sort_values(ascending=False)
)

# 3C. Plot: (C1) Pie of all 144 Items by Yield
figC1, axC1 = plt.subplots(figsize=(10, 10))
colors_items_C = plt.cm.tab20b(np.linspace(0, 1, len(yield_per_item)))
wedgesC1, textsC1, autotextsC1 = axC1.pie(
    yield_per_item.values,
    labels=yield_per_item.index,
    autopct='%1.1f%%',
    colors=colors_items_C,
    startangle=90,
    wedgeprops=dict(edgecolor='white', linewidth=0.5)
)
for txt in textsC1:
    txt.set_fontsize(4)
for atxt in autotextsC1:
    atxt.set_fontsize(3)
axC1.set_title('Total Yield by Crop Item (all 144)', fontsize=14, pad=20)
plt.tight_layout()

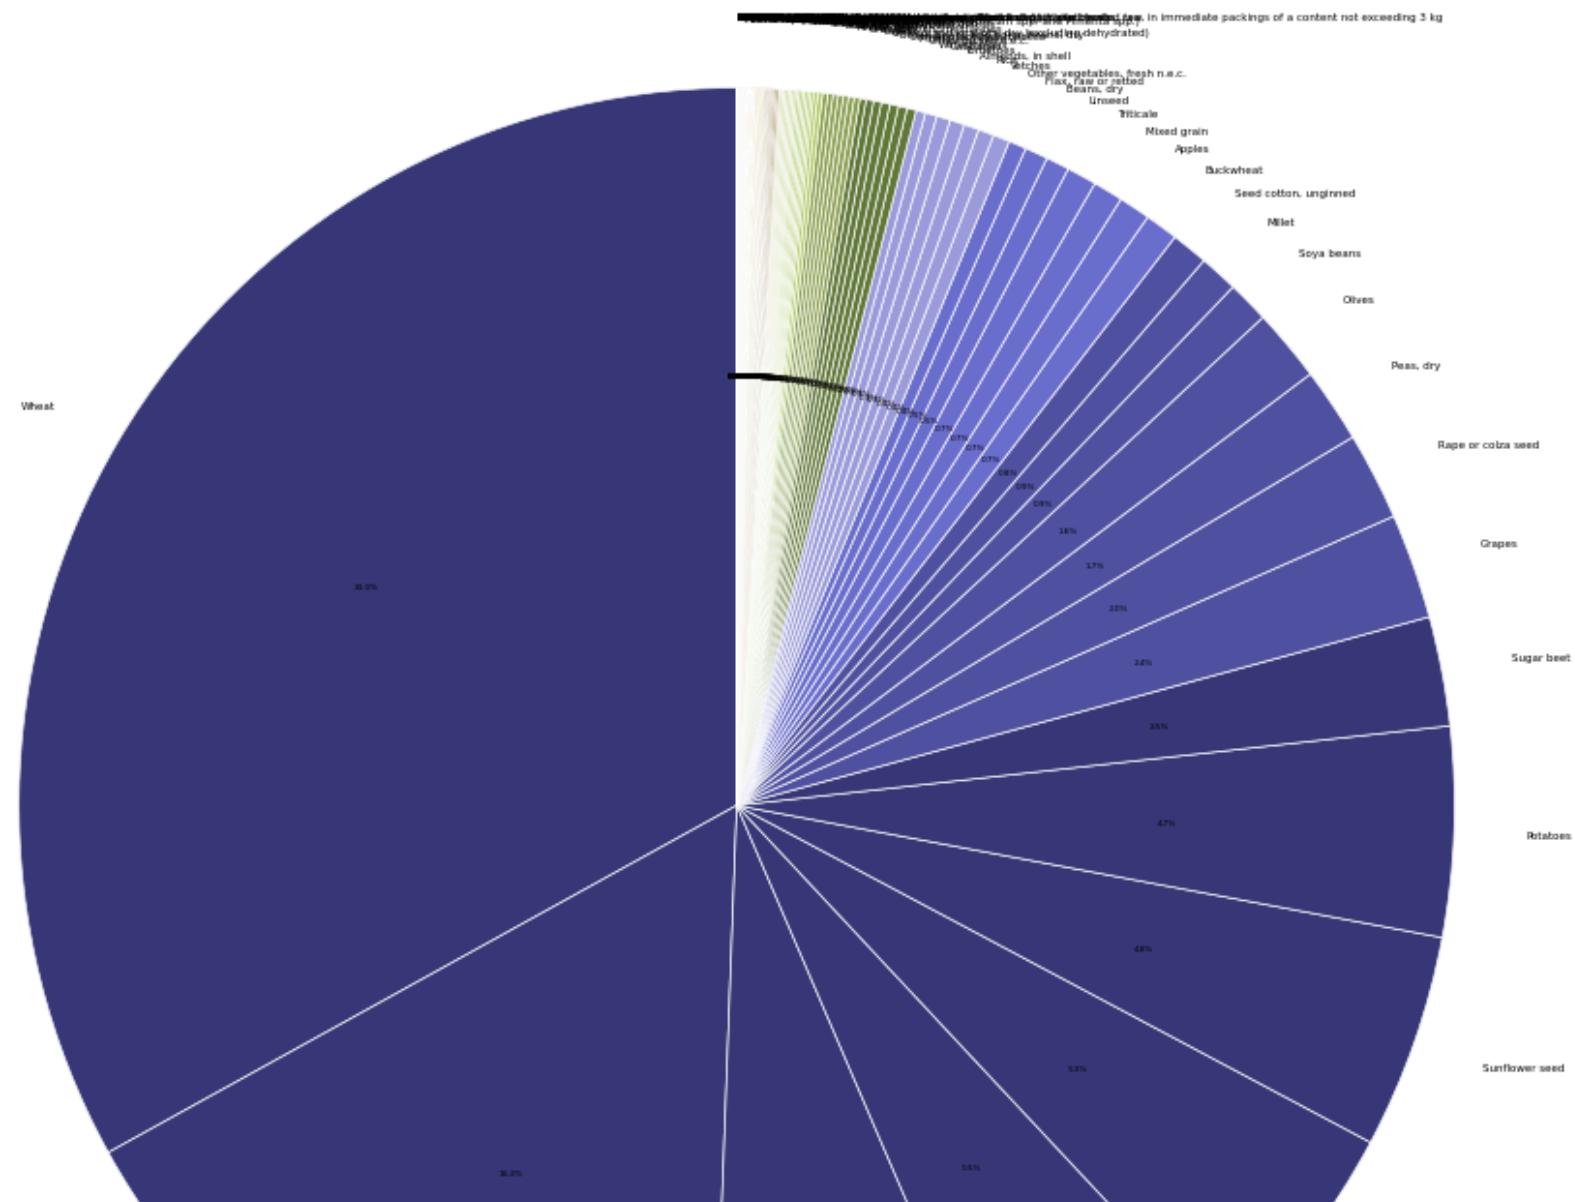
# 4C. Plot: (C2) Pie of Groups by Yield (Legend on right)
figC2, axC2 = plt.subplots(figsize=(12, 8))
groups_C = yield_by_group.index

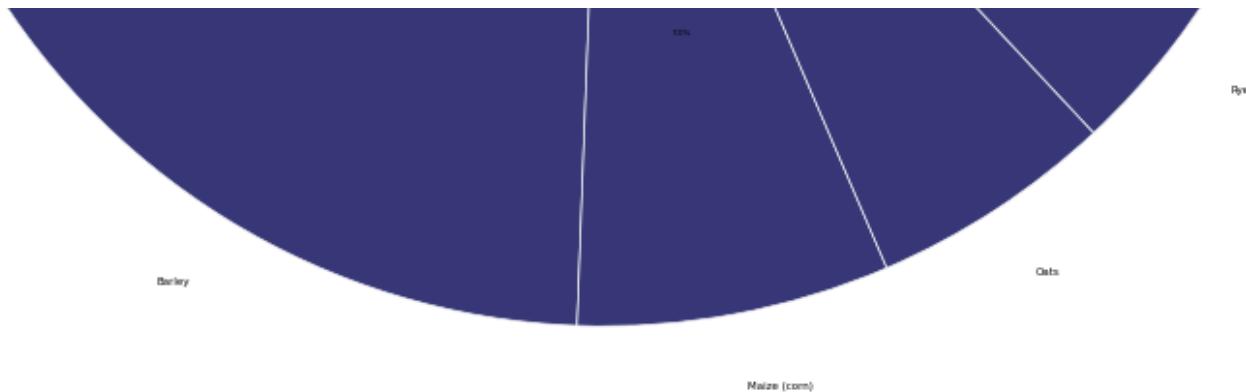
```

```
values_C = yield_by_group.values
total_C = values_C.sum()
pct_C = values_C / total_C * 100
legend_labels_C = [f"{grp}: {p:.3f}%" for grp, p in zip(groups_C, pct_C)]
cmap_C = plt.get_cmap('tab20')
colors_groups_C = cmap_C.colors[: len(groups_C)]
wedgesC2, _ = axC2.pie(
    values_C,
    labels=None,
    startangle=90,
    colors=colors_groups_C,
    wedgeprops=dict(edgecolor='white', linewidth=1)
)
centreC2 = plt.Circle((0, 0), 0.45, color='white', linewidth=0)
axC2.add_artist(centreC2)
axC2.set_title('Total Yield by Category (grouped)', fontsize=16, pad=20)
axC2.legend(
    wedgesC2,
    legend_labels_C,
    title="Category - % of Total Yield",
    loc="center left",
    bbox_to_anchor=(1, 0, 0.3, 1),
    fontsize=10,
    title_fontsize=12
)
plt.tight_layout()

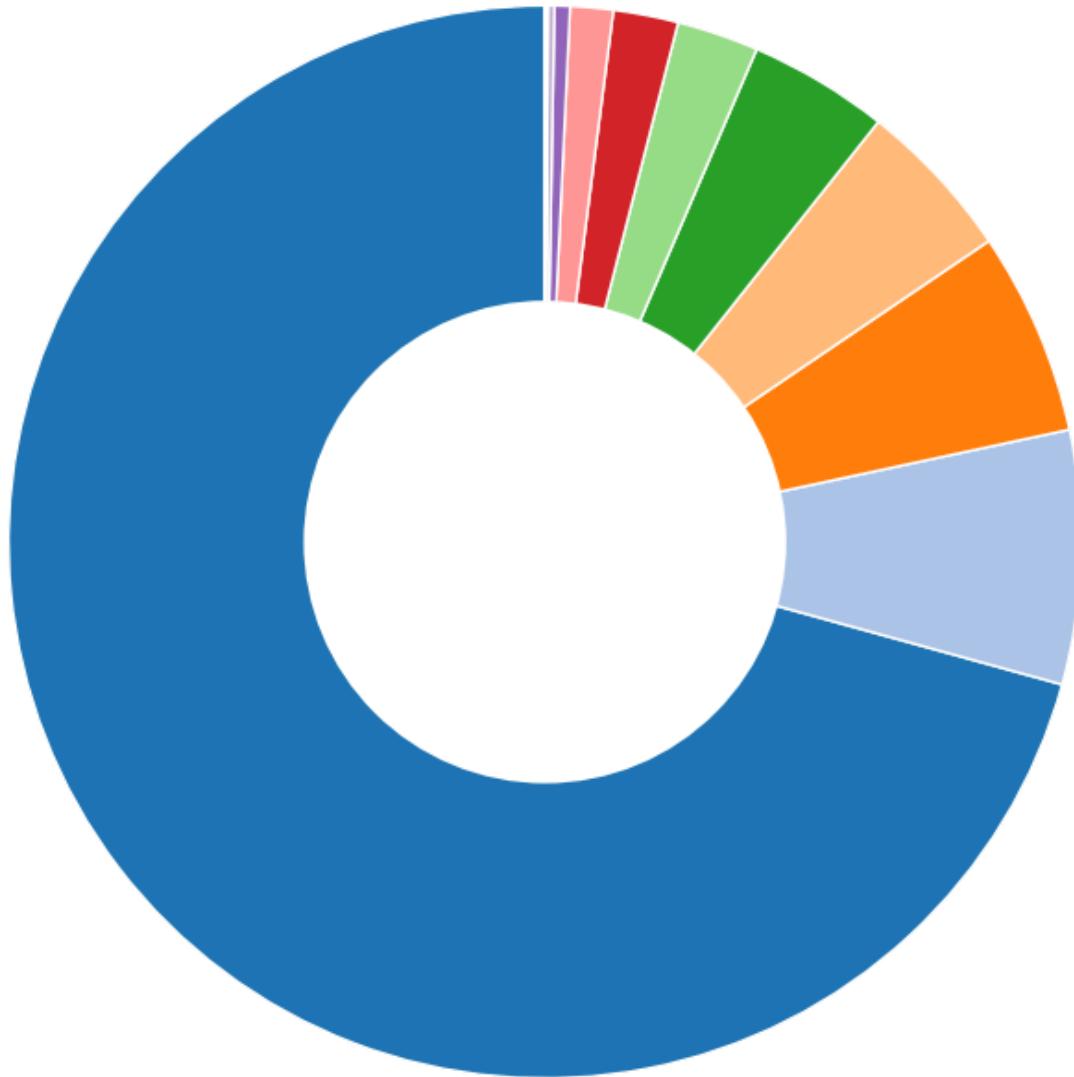
# -----
# DISPLAY ALL FIGURES
# -----
plt.show()
```

Total Area Harvested by Crop Item (all 144)



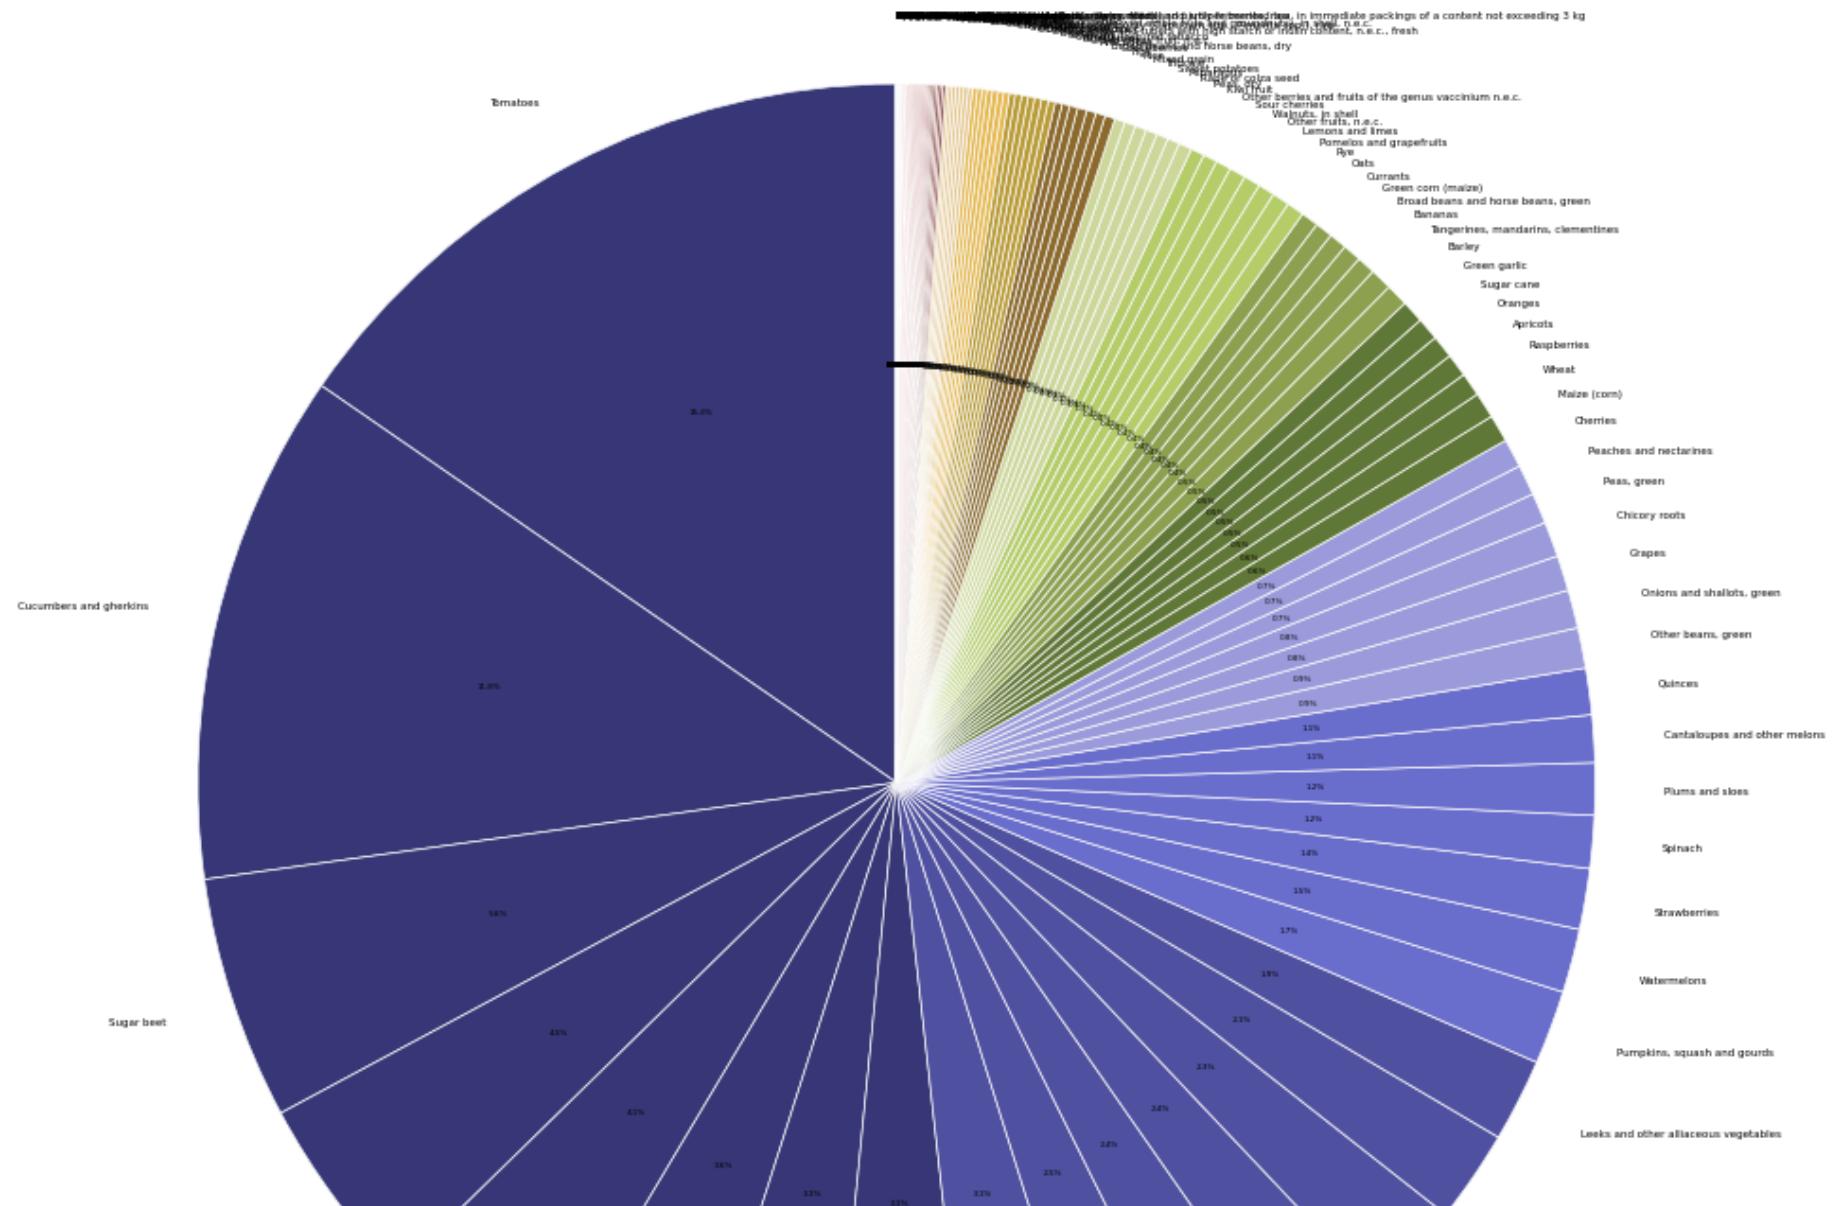


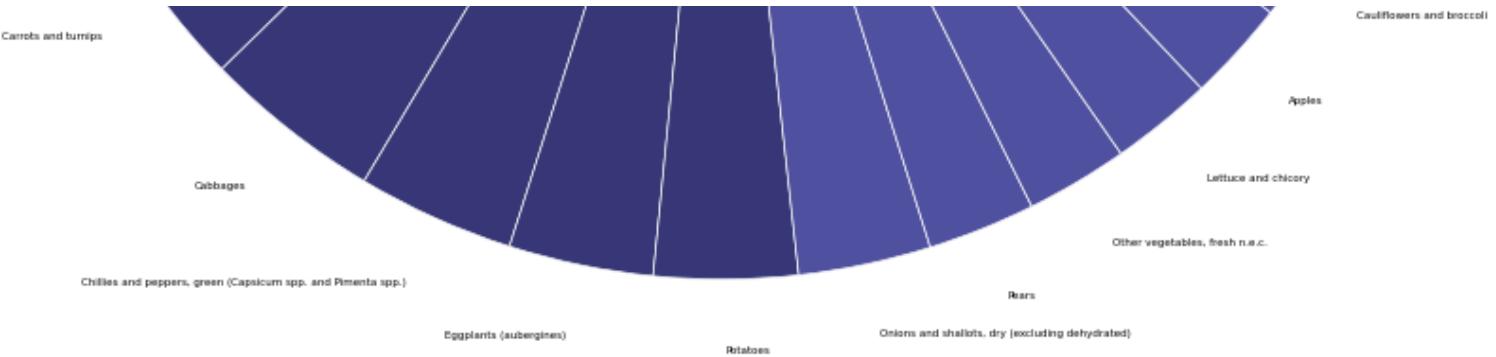
Total Area Harvested by Category (grouped)



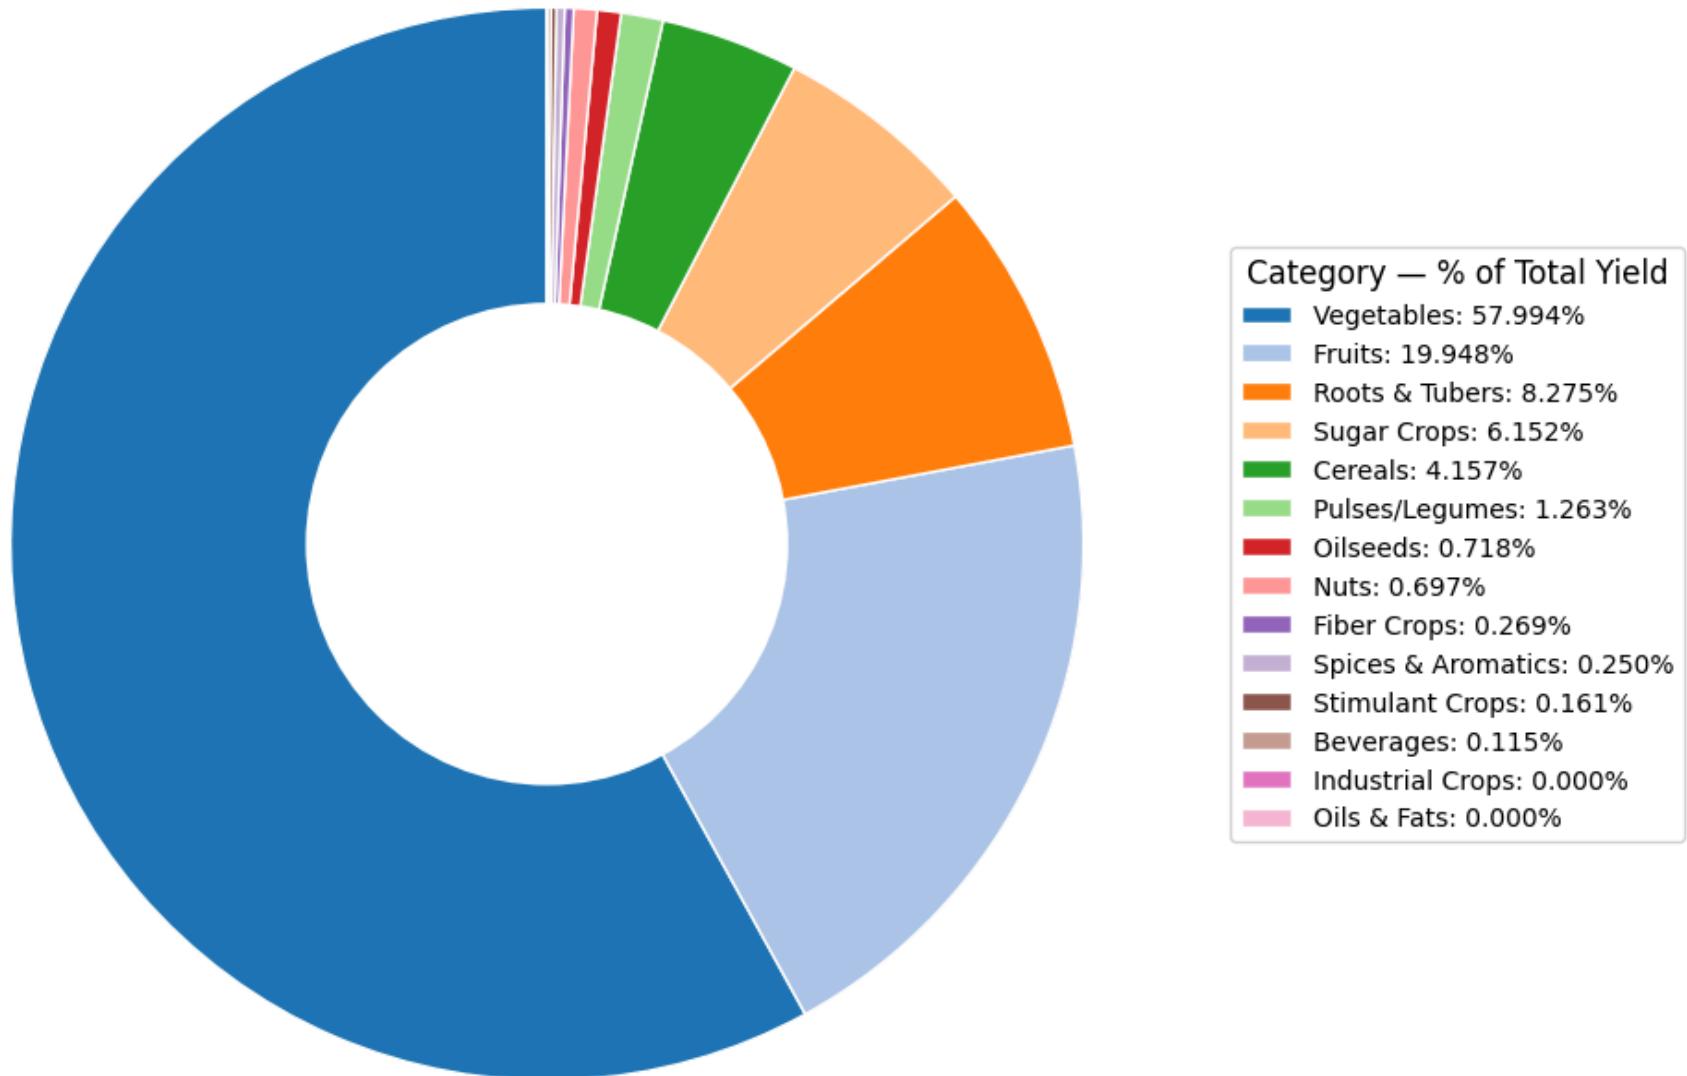
Category	% of Total Area
Cereals	70.718%
Oilseeds	7.643%
Fruits	6.125%
Roots & Tubers	4.840%
Pulses/Legumes	4.217%
Sugar Crops	2.464%
Vegetables	1.958%
Fiber Crops	1.292%
Nuts	0.465%
Stimulant Crops	0.195%
Spices & Aromatics	0.045%
Beverages	0.040%
Industrial Crops	0.000%
Oils & Fats	0.000%

Total Yield by Crop Item (all 144)





Total Yield by Category (grouped)



Now that we have established that given the number of crops in our dataset we kind of have to group them i have chose to alter the dataset such that these

```
In [76]: # 1) Copy df_merged so the original remains unchanged
df_temp = df_merged.copy()

# 2) Add the 'Group' column by mapping each 'Item' to its category
df_temp['Group'] = df_temp['Item'].map(crop_to_group)

# 3) Define a helper to pick the first non-null value in a series
def first_non_null(series):
    non_null = series.dropna()
    return non_null.iloc[0] if not non_null.empty else np.nan

# 4) Group by both 'Group' and 'Year' to retain year information
df_grouped = (
    df_temp
    .groupby(['Group', 'Year'])
    .agg({
        'Area_harvested':      'sum',
        'Unit_Area_harvested': first_non_null,
        'Production':          'sum',
        'Unit_Production':     first_non_null,
        'Yield':                'sum',
        'Unit_Yield':           first_non_null
    })
    .reset_index()
)

# 5) (Optional) To display Large numbers without scientific notation:
pd.set_option('display.float_format', '{:.0f}'.format)

# 6) View the result
print(df_grouped)
```

	Group	Year	Area_harvested	Unit_Area_harvested	Production	\
0	Beverages	1961	105,975	ha	39,687,081	
1	Beverages	1962	110,761	ha	46,654,232	
2	Beverages	1963	117,763	ha	44,456,323	
3	Beverages	1964	121,524	ha	49,621,382	
4	Beverages	1965	123,504	ha	50,226,655	
..
870	Vegetables	2019	2,878,939	ha	76,594,172	
871	Vegetables	2020	2,907,378	ha	77,633,080	
872	Vegetables	2021	2,907,543	ha	79,882,333	
873	Vegetables	2022	2,675,089	ha	71,661,490	
874	Vegetables	2023	2,686,924	ha	72,559,786	
	Unit_Production		Yield	Unit_Yield		
0	t	15,032	kg/ha			
1	t	15,989	kg/ha			
2	t	17,234	kg/ha			
3	t	18,440	kg/ha			
4	t	18,082	kg/ha			
..		
870	t	19,465,401	kg/ha			
871	t	19,975,897	kg/ha			
872	t	20,327,850	kg/ha			
873	t	19,863,159	kg/ha			
874	t	19,937,441	kg/ha			

[875 rows x 8 columns]

```
In [80]: # 2) Aggregate across years so that each Group has a single total for each metric
df_summary = df_grouped.groupby('Group').agg({
    'Area_harvested': 'sum',
    'Production':      'sum',
    'Yield':           'sum'
}).reset_index()

# (Optional) Sort by Production so bars appear in descending order of production
df_summary = df_summary.sort_values('Production', ascending=False)

# 3) Create a figure with three subplots side by side
fig, axes = plt.subplots(1, 3, figsize=(18, 6), sharex=True)

# 3A) Bar chart: Total Area_harvested by Group (no units in ylabel)
```

```

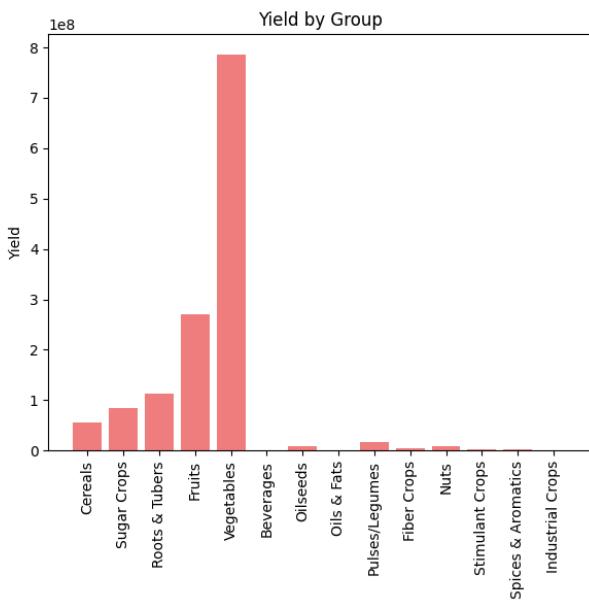
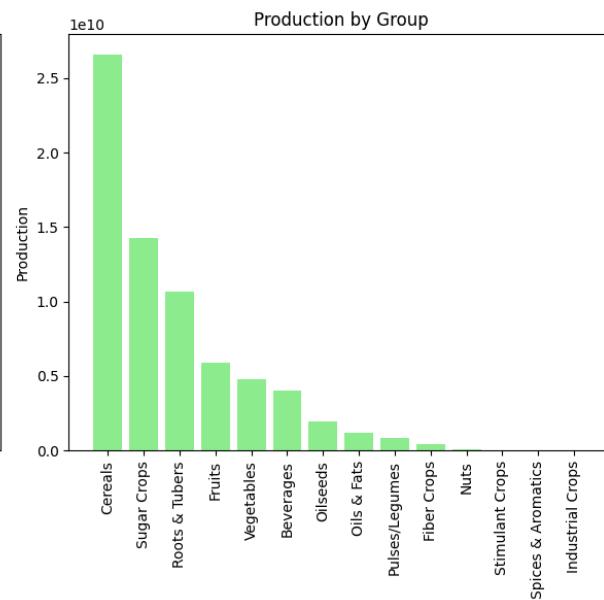
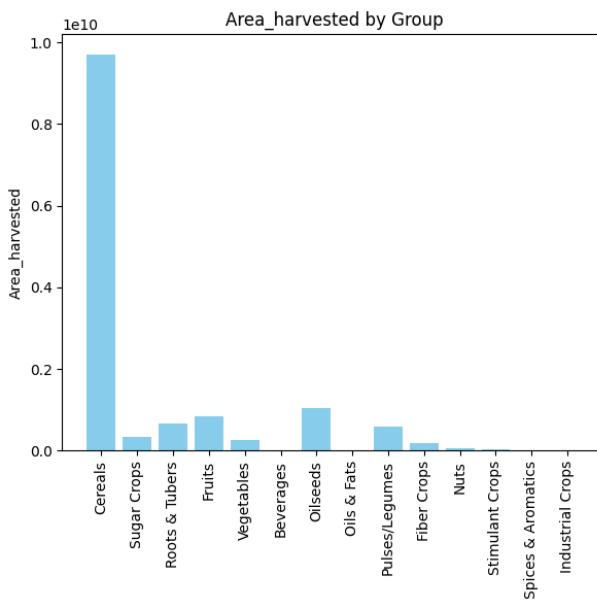
axes[0].bar(df_summary['Group'], df_summary['Area_harvested'], color='skyblue')
axes[0].set_title('Area_harvested by Group')
axes[0].set_ylabel('Area_harvested')
axes[0].tick_params(axis='x', rotation=90)

# 3B) Bar chart: Total Production by Group (no units in ylabel)
axes[1].bar(df_summary['Group'], df_summary['Production'], color='lightgreen')
axes[1].set_title('Production by Group')
axes[1].set_ylabel('Production')
axes[1].tick_params(axis='x', rotation=90)

# 3C) Bar chart: Total Yield by Group (no units in ylabel)
axes[2].bar(df_summary['Group'], df_summary['Yield'], color='lightcoral')
axes[2].set_title('Yield by Group')
axes[2].set_ylabel('Yield')
axes[2].tick_params(axis='x', rotation=90)

# 4) Tidy up layout so labels aren't cut off
plt.tight_layout()
plt.show()

```



In [97]: # 2) Aggregate across years so that each Group has a single total for each metric
df_summary = df_grouped.groupby('Group').agg({

```
'Area_harvested': 'sum',
'Production':      'sum',
'Yield':           'sum'
}).reset_index()

# (Optional) Sort by Production so bars appear in descending order of production
df_summary = df_summary.sort_values('Production', ascending=False)

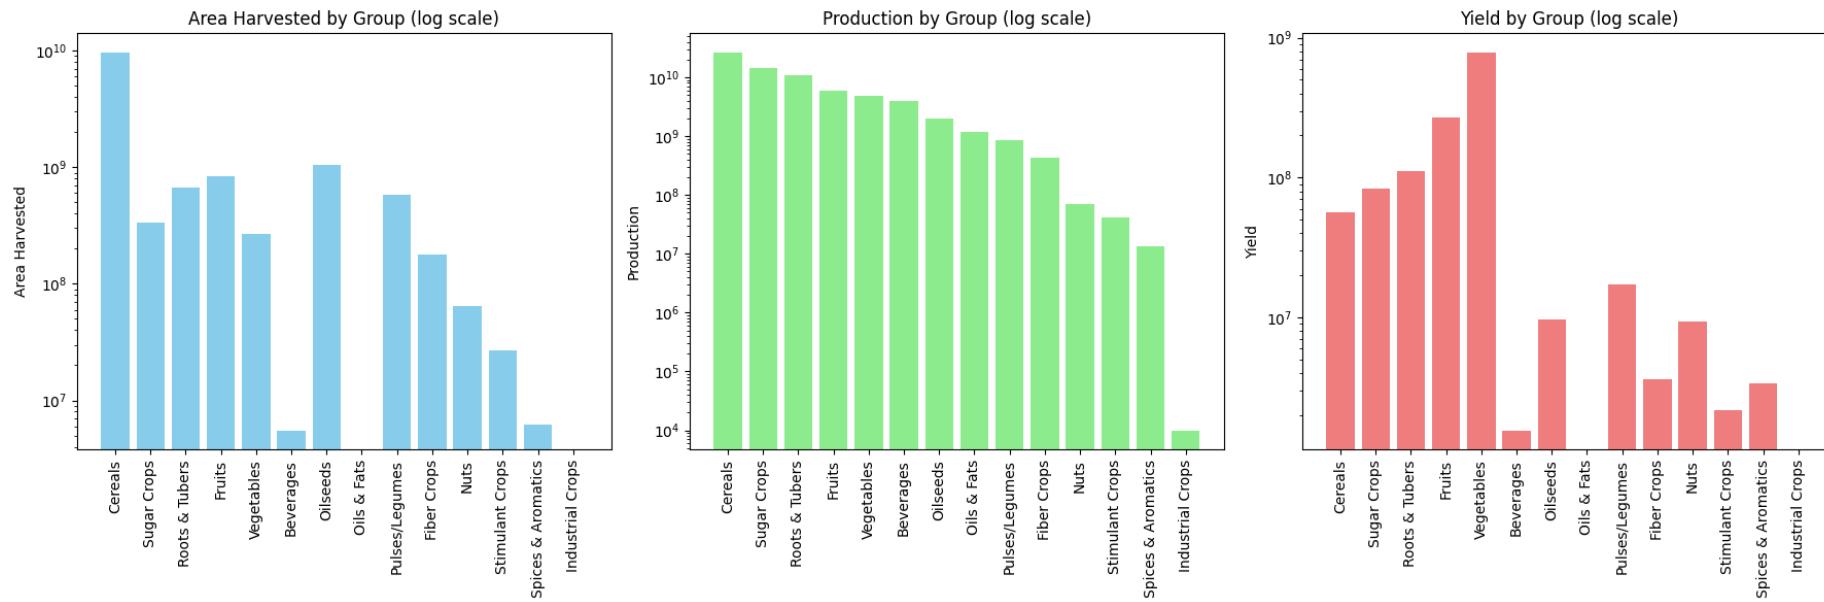
# 3) Create a figure with three subplots side by side
fig, axes = plt.subplots(1, 3, figsize=(18, 6), sharex=False)

# 3A) Bar chart: Total Area_harvested by Group (log scale)
axes[0].bar(df_summary['Group'], df_summary['Area_harvested'], color='skyblue')
axes[0].set_title('Area Harvested by Group (log scale)')
axes[0].set_ylabel('Area Harvested')
axes[0].tick_params(axis='x', rotation=90)
axes[0].set_yscale('log')

# 3B) Bar chart: Total Production by Group (log scale)
axes[1].bar(df_summary['Group'], df_summary['Production'], color='lightgreen')
axes[1].set_title('Production by Group (log scale)')
axes[1].set_ylabel('Production')
axes[1].tick_params(axis='x', rotation=90)
axes[1].set_yscale('log')

# 3C) Bar chart: Total Yield by Group (log scale)
axes[2].bar(df_summary['Group'], df_summary['Yield'], color='lightcoral')
axes[2].set_title('Yield by Group (log scale)')
axes[2].set_ylabel('Yield')
axes[2].tick_params(axis='x', rotation=90)
axes[2].set_yscale('log')

# 4) Tidy up Layout so Labels aren't cut off
plt.tight_layout()
plt.show()
```



```
In [90]: # Create pivot tables for each metric
pivot_area = df_grouped.pivot(index='Group', columns='Year', values='Area_harvested').fillna(0)
pivot_production = df_grouped.pivot(index='Group', columns='Year', values='Production').fillna(0)
pivot_yield = df_grouped.pivot(index='Group', columns='Year', values='Yield').fillna(0)

# Plotting heatmaps using matplotlib
fig, axes = plt.subplots(3, 1, figsize=(12, 18), constrained_layout=True)

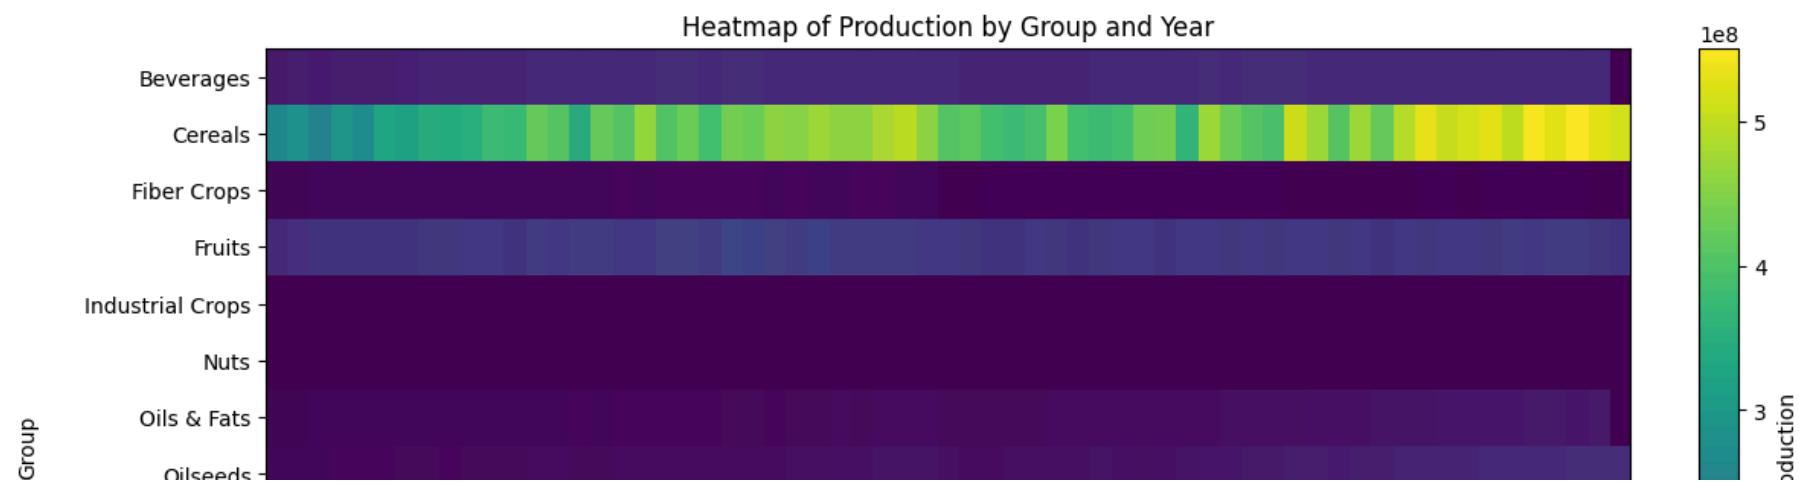
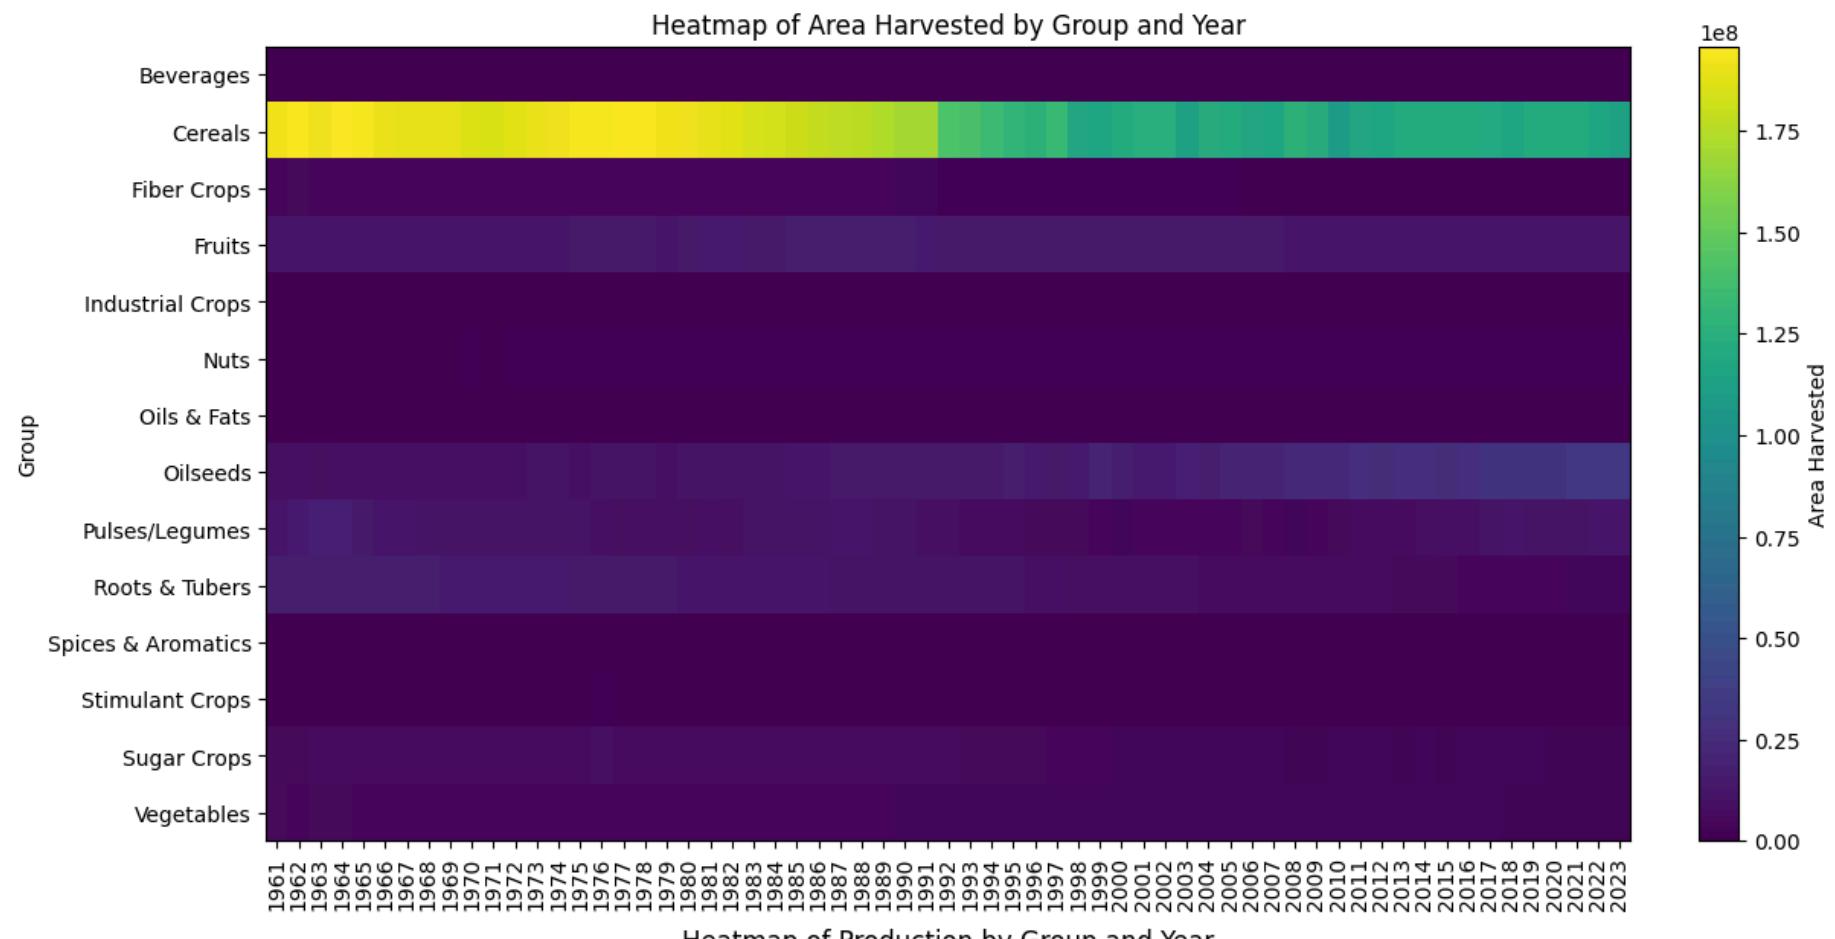
# 1) Heatmap for Area harvested
im1 = axes[0].imshow(pivot_area, aspect='auto', cmap='viridis')
axes[0].set_title('Heatmap of Area Harvested by Group and Year')
axes[0].set_ylabel('Group')
axes[0].set_xticks(np.arange(len(pivot_area.columns)))
axes[0].set_yticks(np.arange(len(pivot_area.index)))
axes[0].set_xticklabels(pivot_area.columns, rotation=90)
axes[0].set_yticklabels(pivot_area.index)
fig.colorbar(im1, ax=axes[0], orientation='vertical', label='Area Harvested')

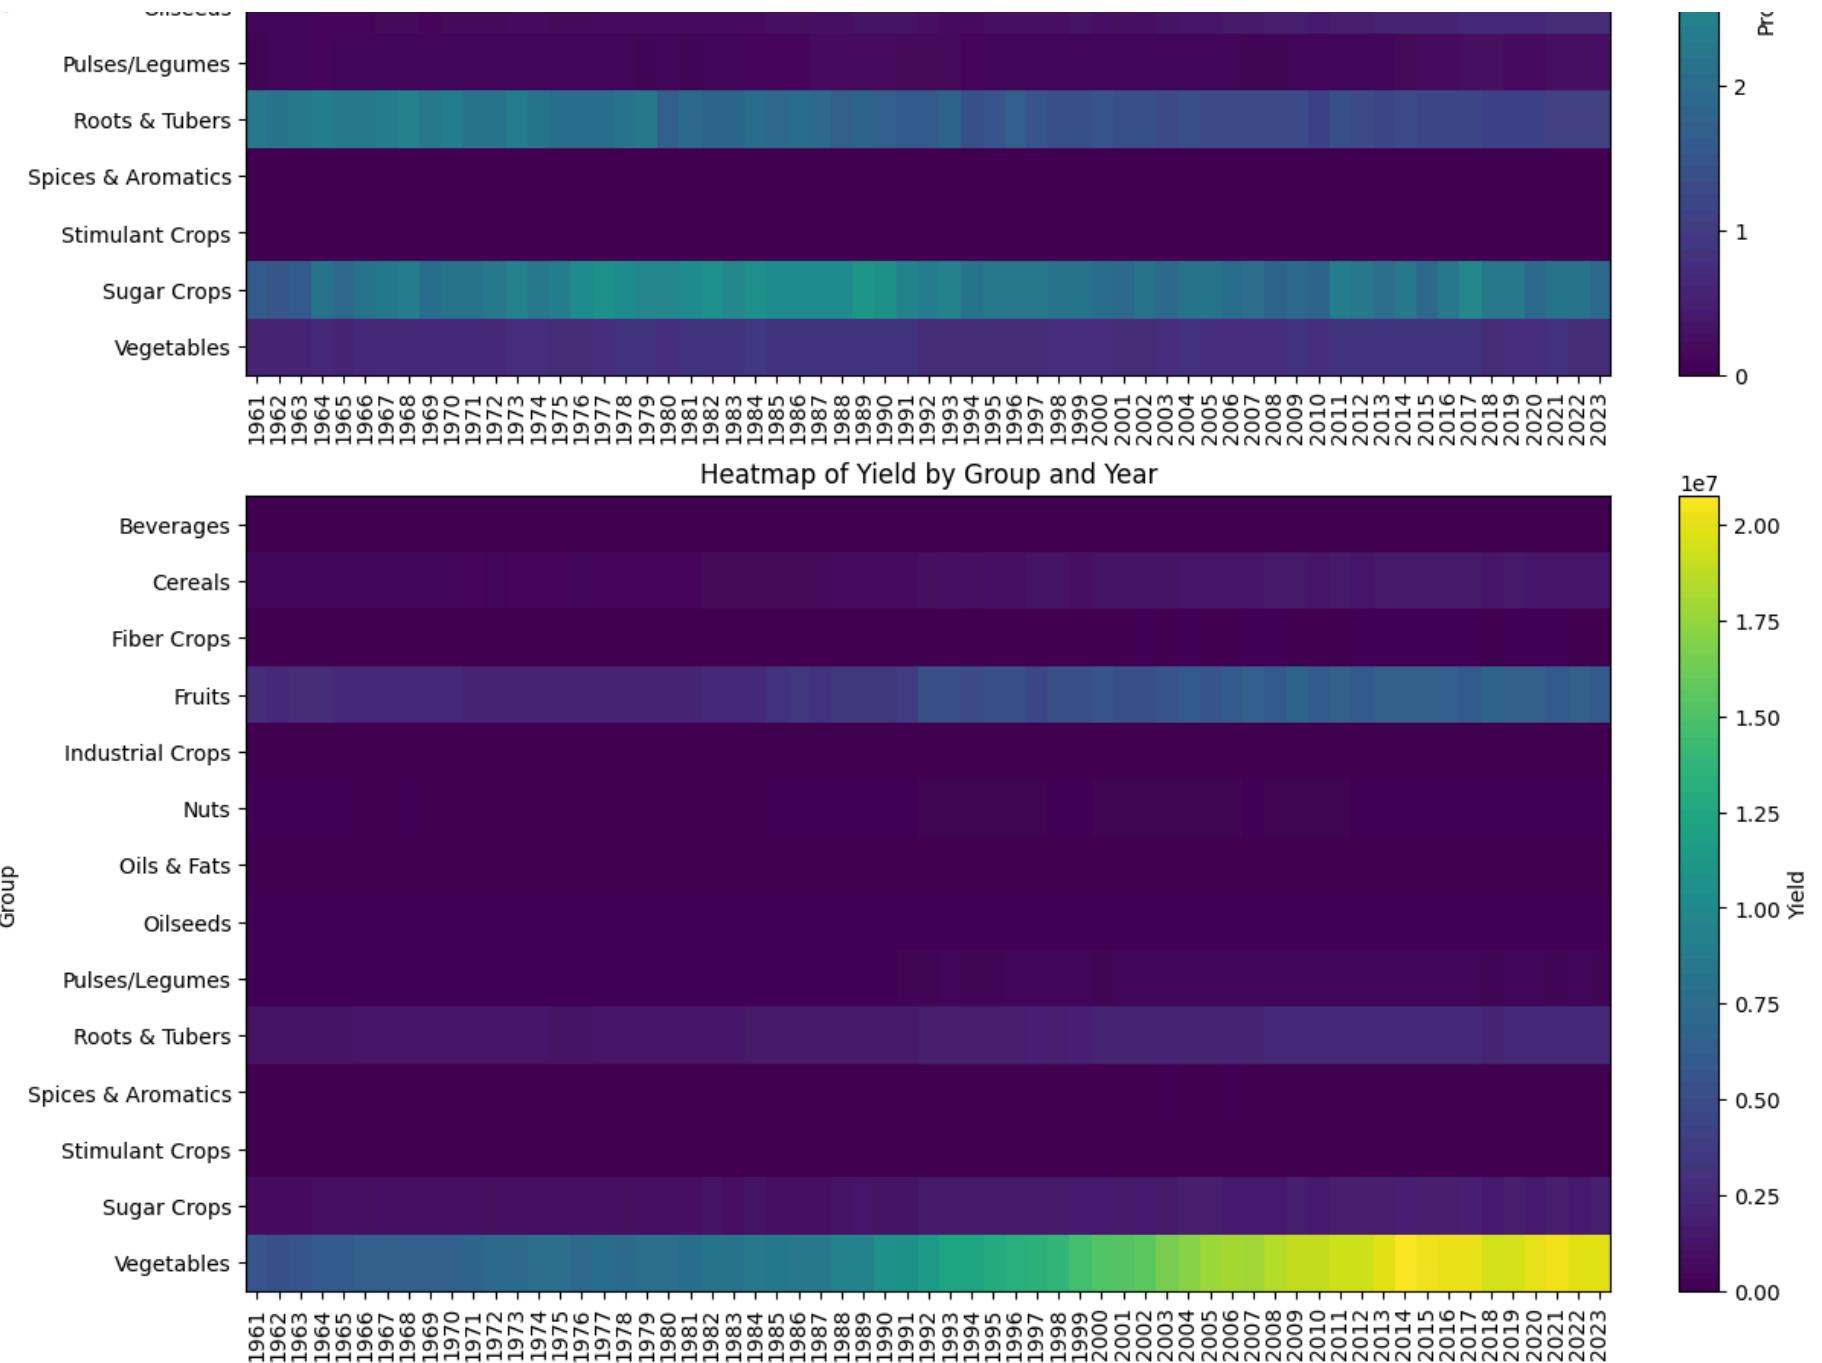
# 2) Heatmap for Production
im2 = axes[1].imshow(pivot_production, aspect='auto', cmap='viridis')
axes[1].set_title('Heatmap of Production by Group and Year')
axes[1].set_ylabel('Group')
```

```
axes[1].set_xticks(np.arange(len(pivot_production.columns)))
axes[1].set_yticks(np.arange(len(pivot_production.index)))
axes[1].set_xticklabels(pivot_production.columns, rotation=90)
axes[1].set_yticklabels(pivot_production.index)
fig.colorbar(im2, ax=axes[1], orientation='vertical', label='Production')

# 3) Heatmap for Yield
im3 = axes[2].imshow(pivot_yield, aspect='auto', cmap='viridis')
axes[2].set_title('Heatmap of Yield by Group and Year')
axes[2].set_ylabel('Group')
axes[2].set_xticks(np.arange(len(pivot_yield.columns)))
axes[2].set_yticks(np.arange(len(pivot_yield.index)))
axes[2].set_xticklabels(pivot_yield.columns, rotation=90)
axes[2].set_yticklabels(pivot_yield.index)
fig.colorbar(im3, ax=axes[2], orientation='vertical', label='Yield')

plt.show()
```





In [89]:

```
# Create pivot tables for each metric
pivot_area = df_grouped.pivot(index='Group', columns='Year', values='Area_harvested').fillna(0)
pivot_production = df_grouped.pivot(index='Group', columns='Year', values='Production').fillna(0)
pivot_yield = df_grouped.pivot(index='Group', columns='Year', values='Yield').fillna(0)

# Determine nonzero minima for log scaling
area_min = pivot_area[pivot_area > 0].min().min()
prod_min = pivot_production[pivot_production > 0].min().min()
yield_min = pivot_yield[pivot_yield > 0].min().min()

# Plotting heatmaps using matplotlib with Log normalization
fig, axes = plt.subplots(3, 1, figsize=(12, 18), constrained_layout=True)

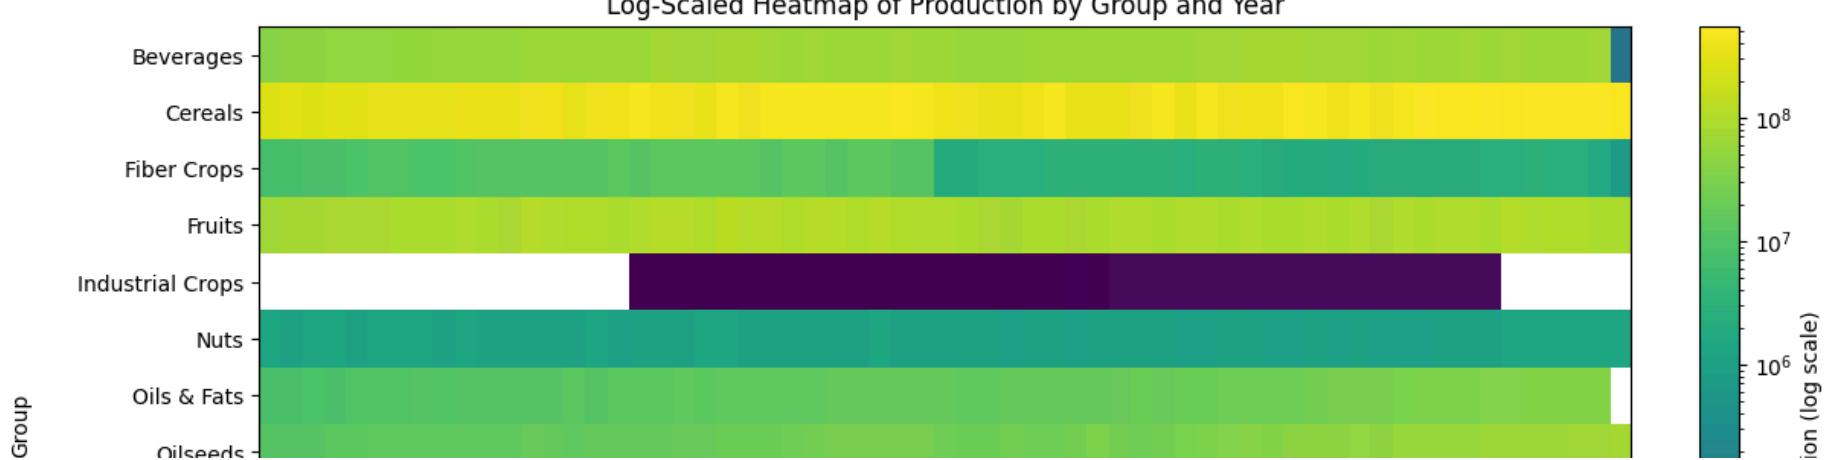
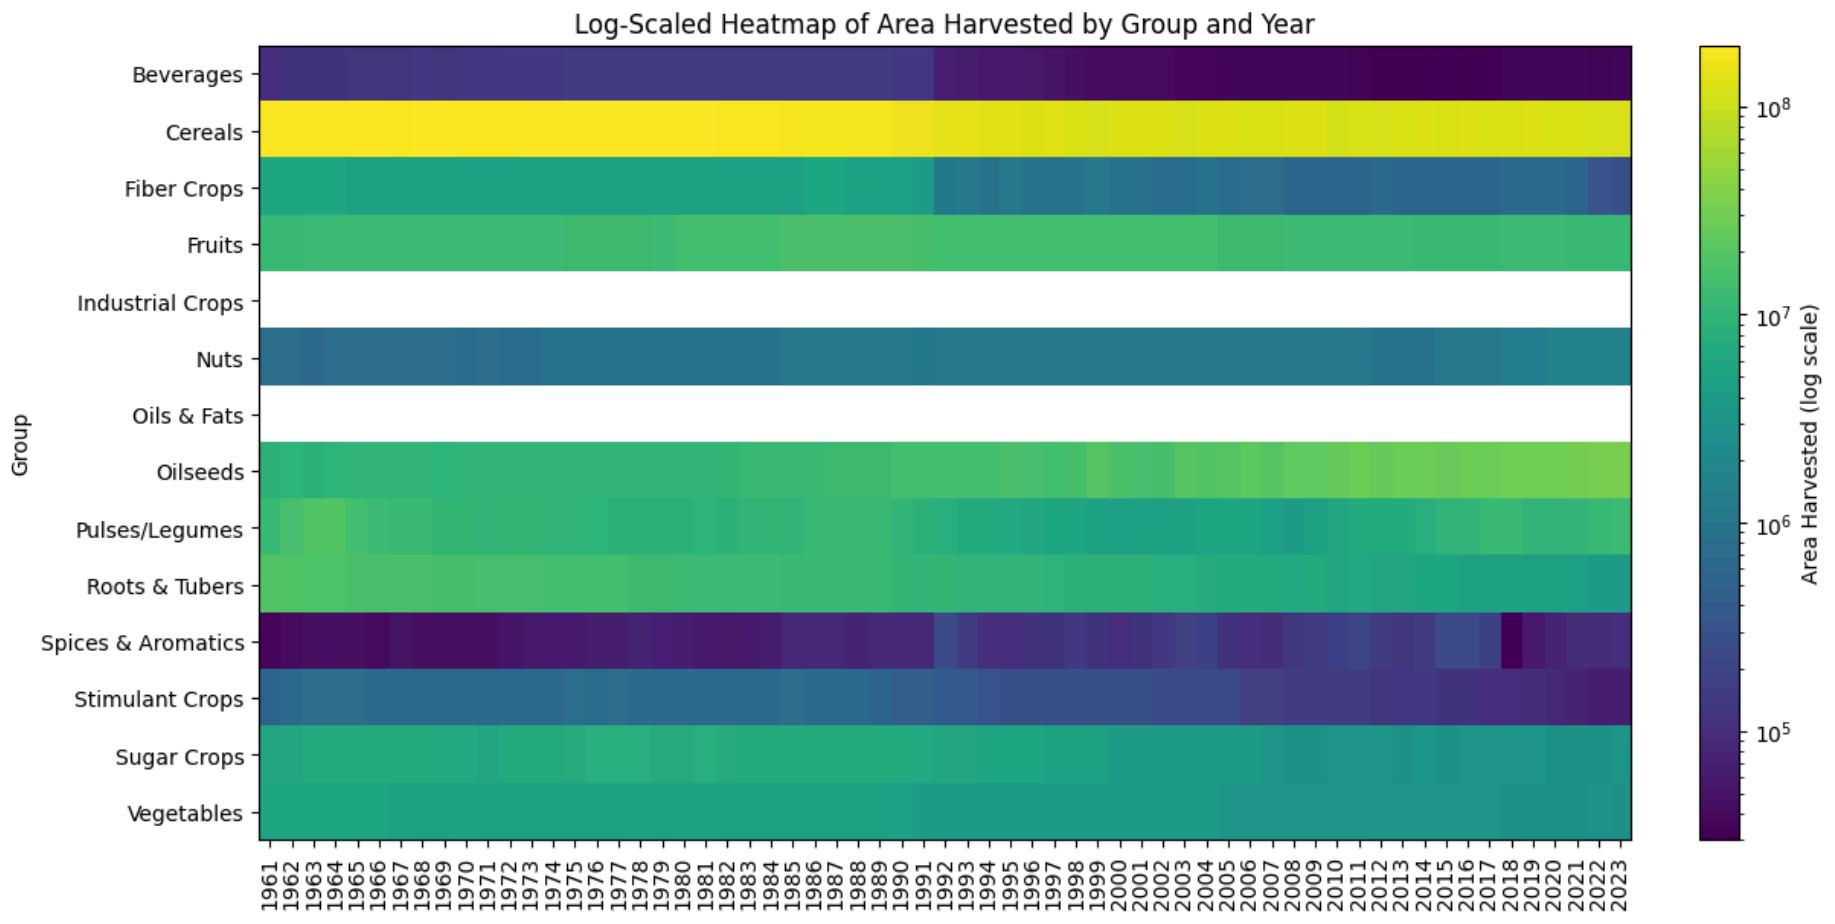
# 1) Heatmap for Area_harvested (Log scale)
im1 = axes[0].imshow(
    pivot_area,
    aspect='auto',
    cmap='viridis',
    norm=LogNorm(vmin=area_min, vmax=pivot_area.max().max())
)
axes[0].set_title('Log-Scaled Heatmap of Area Harvested by Group and Year')
axes[0].set_ylabel('Group')
axes[0].set_xticks(np.arange(len(pivot_area.columns)))
axes[0].set_yticks(np.arange(len(pivot_area.index)))
axes[0].set_xticklabels(pivot_area.columns, rotation=90)
axes[0].set_yticklabels(pivot_area.index)
fig.colorbar(im1, ax=axes[0], orientation='vertical', label='Area Harvested (log scale)')

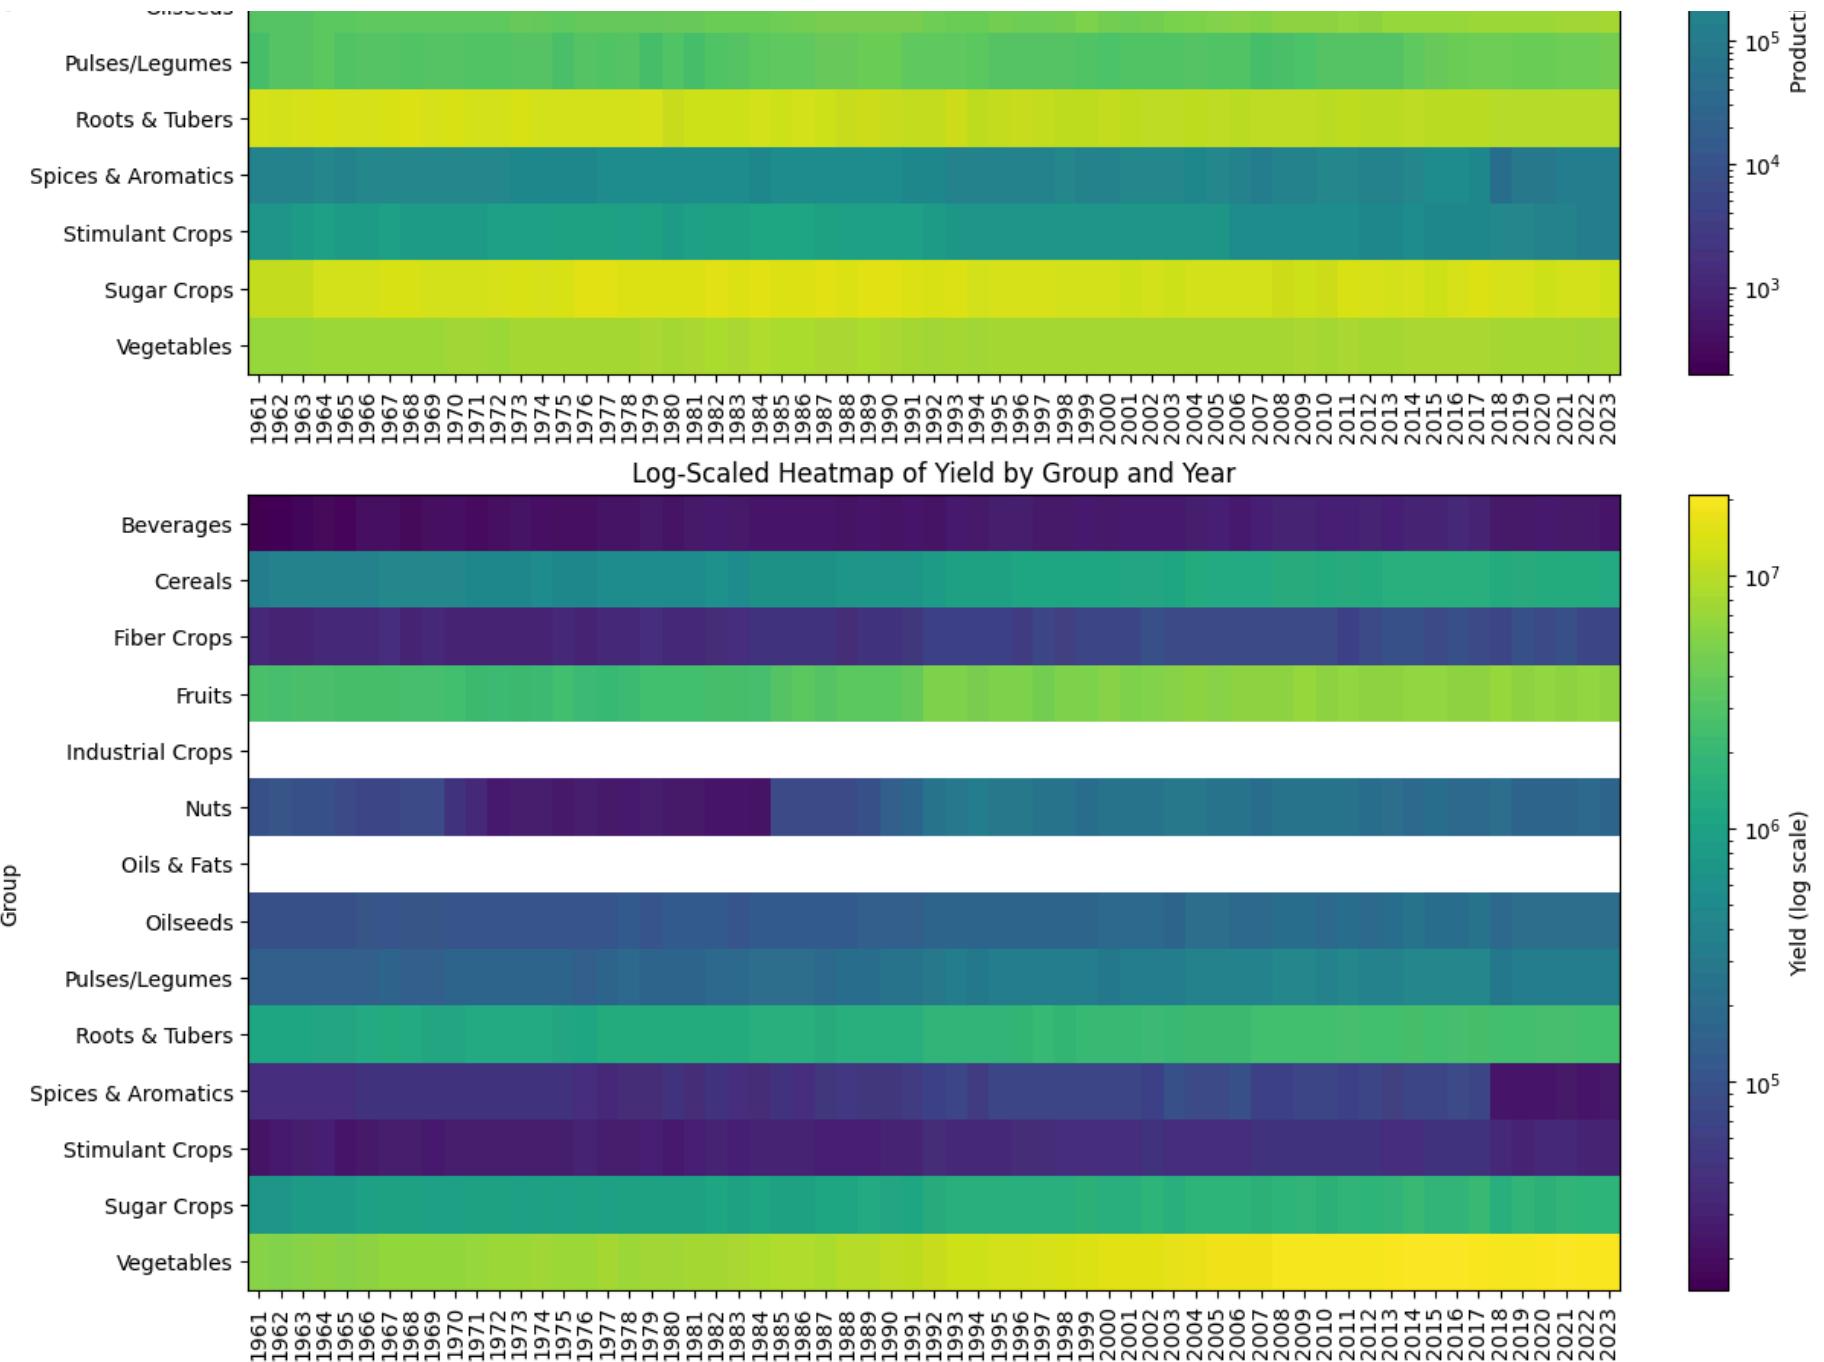
# 2) Heatmap for Production (Log scale)
im2 = axes[1].imshow(
    pivot_production,
    aspect='auto',
    cmap='viridis',
    norm=LogNorm(vmin=prod_min, vmax=pivot_production.max().max())
)
axes[1].set_title('Log-Scaled Heatmap of Production by Group and Year')
axes[1].set_ylabel('Group')
axes[1].set_xticks(np.arange(len(pivot_production.columns)))
axes[1].set_yticks(np.arange(len(pivot_production.index)))
axes[1].set_xticklabels(pivot_production.columns, rotation=90)
axes[1].set_yticklabels(pivot_production.index)
```

```
fig.colorbar(im2, ax=axes[1], orientation='vertical', label='Production (log scale)')

# 3) Heatmap for Yield (log scale)
im3 = axes[2].imshow(
    pivot_yield,
    aspect='auto',
    cmap='viridis',
    norm=LogNorm(vmin=yield_min, vmax=pivot_yield.max().max()))
)
axes[2].set_title('Log-Scaled Heatmap of Yield by Group and Year')
axes[2].set_ylabel('Group')
axes[2].set_xticks(np.arange(len(pivot_yield.columns)))
axes[2].set_yticks(np.arange(len(pivot_yield.index)))
axes[2].set_xticklabels(pivot_yield.columns, rotation=90)
axes[2].set_yticklabels(pivot_yield.index)
fig.colorbar(im3, ax=axes[2], orientation='vertical', label='Yield (log scale)')

plt.show()
```





In [95]:

```
df_summary = (
    df_grouped
    .groupby('Group')
    .agg({
        'Area_harvested': 'sum',
        'Production': 'sum',
        'Yield': 'sum'
    })
    .reset_index()
)

# Sort (e.g. by Production) so the order is consistent in all treemaps:
df_summary = df_summary.sort_values('Production', ascending=False)

# -----
# Now draw three treemaps stacked **vertically**, but with a very large figure
# so that small rectangles get more absolute pixel-space.
# -----


fig, axes = plt.subplots(3, 1, figsize=(14, 24), constrained_layout=True)

# 1) TREEMAP FOR AREA_HARVESTED
df_area = df_summary[df_summary['Area_harvested'] > 0].copy()
sizes_area = df_area['Area_harvested'].values
labels_area = [f"{grp}\n{val:,.0f}" for grp, val in zip(df_area['Group'], df_area['Area_harvested'])]

axes[0].set_title("Treemap of Total Area Harvested by Group", fontsize=16)
squarify.plot(
    sizes=sizes_area,
    label=labels_area,
    ax=axes[0],
    pad=0,                      # no padding → each rectangle uses max space
    color=plt.cm.tab20c.colors[:len(sizes_area)],
    text_kwargs={'fontsize': 10}   # moderately sized text
)
axes[0].axis('off')

# 2) TREEMAP FOR PRODUCTION
df_prod = df_summary[df_summary['Production'] > 0].copy()
sizes_prod = df_prod['Production'].values
```

```
labels_prod = [f"\{grp}\n\{val:,.0f}" for grp, val in zip(df_prod['Group'], df_prod['Production'])]

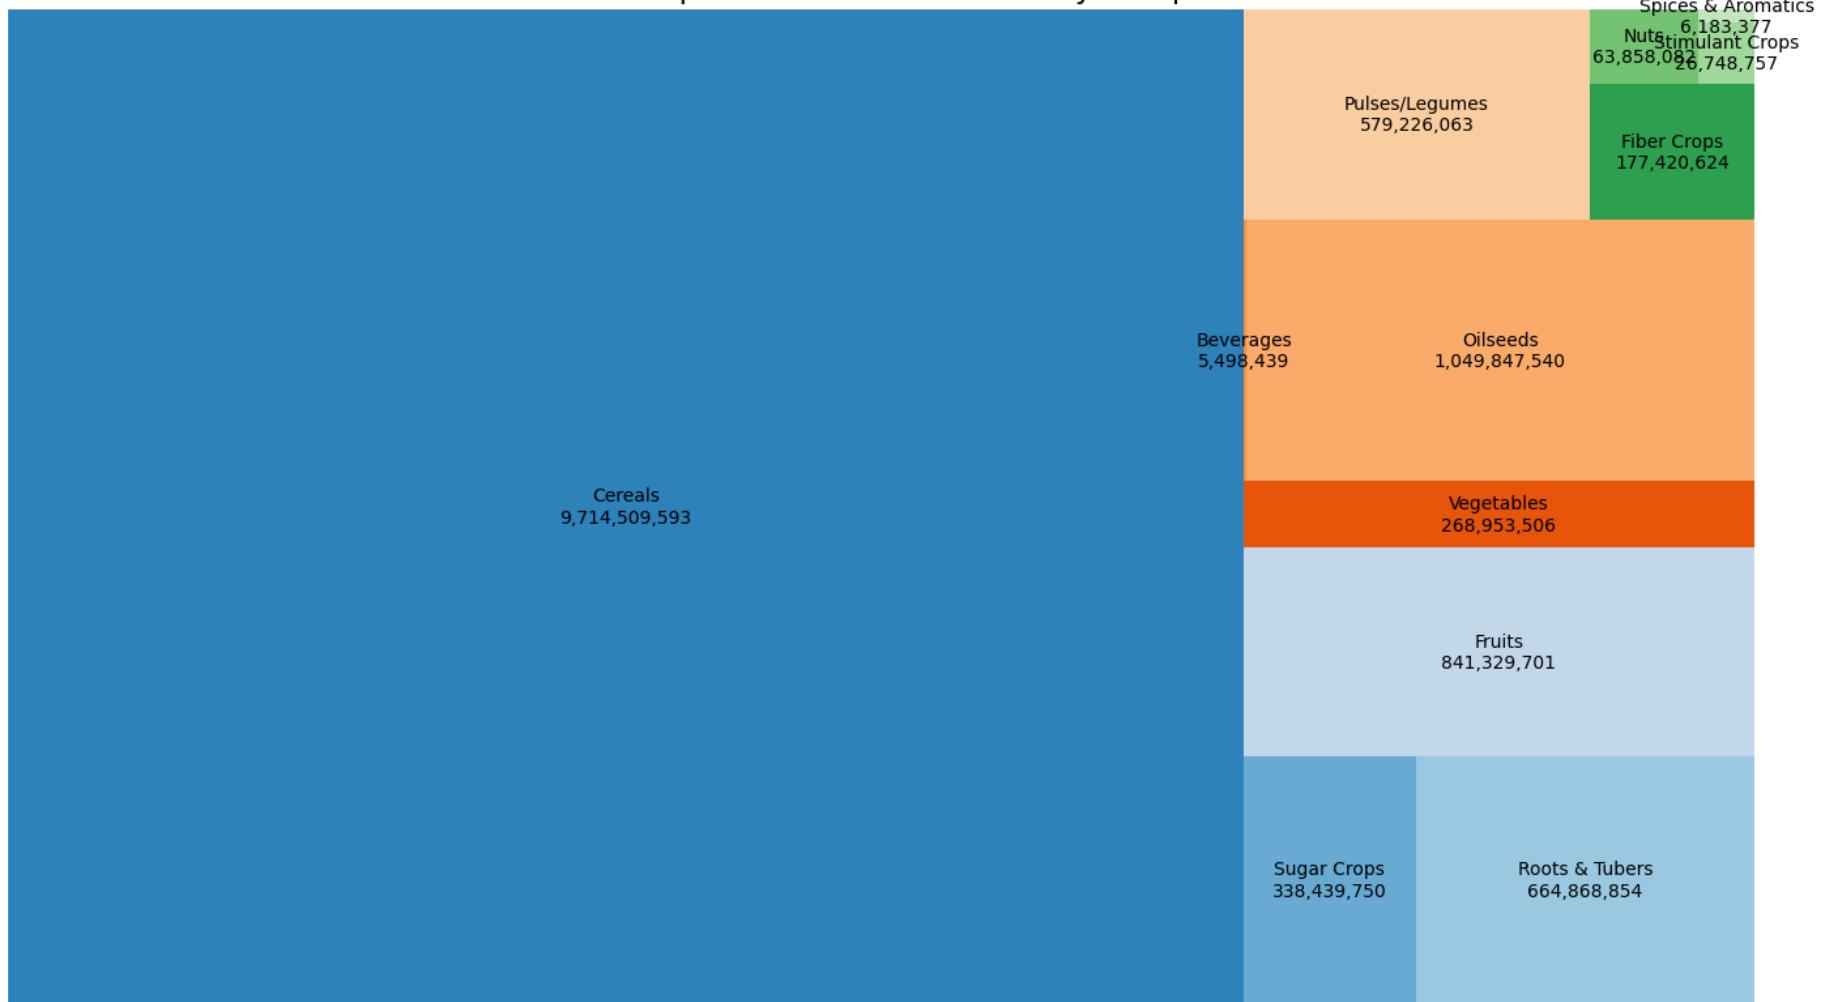
axes[1].set_title("Treemap of Total Production by Group", fontsize=16)
squarify.plot(
    sizes=sizes_prod,
    label=labels_prod,
    ax=axes[1],
    pad=0,
    color=plt.cm.tab20c.colors[:len(sizes_prod)],
    text_kwarg={ 'fontsize': 10}
)
axes[1].axis('off')

# 3) TREEMAP FOR YIELD
df_yield = df_summary[df_summary['Yield'] > 0].copy()
sizes_yield = df_yield['Yield'].values
labels_yield = [f"\{grp}\n\{val:,.0f}" for grp, val in zip(df_yield['Group'], df_yield['Yield'])]

axes[2].set_title("Treemap of Total Yield by Group", fontsize=16)
squarify.plot(
    sizes=sizes_yield,
    label=labels_yield,
    ax=axes[2],
    pad=0,
    color=plt.cm.tab20c.colors[:len(sizes_yield)],
    text_kwarg={ 'fontsize': 10}
)
axes[2].axis('off')

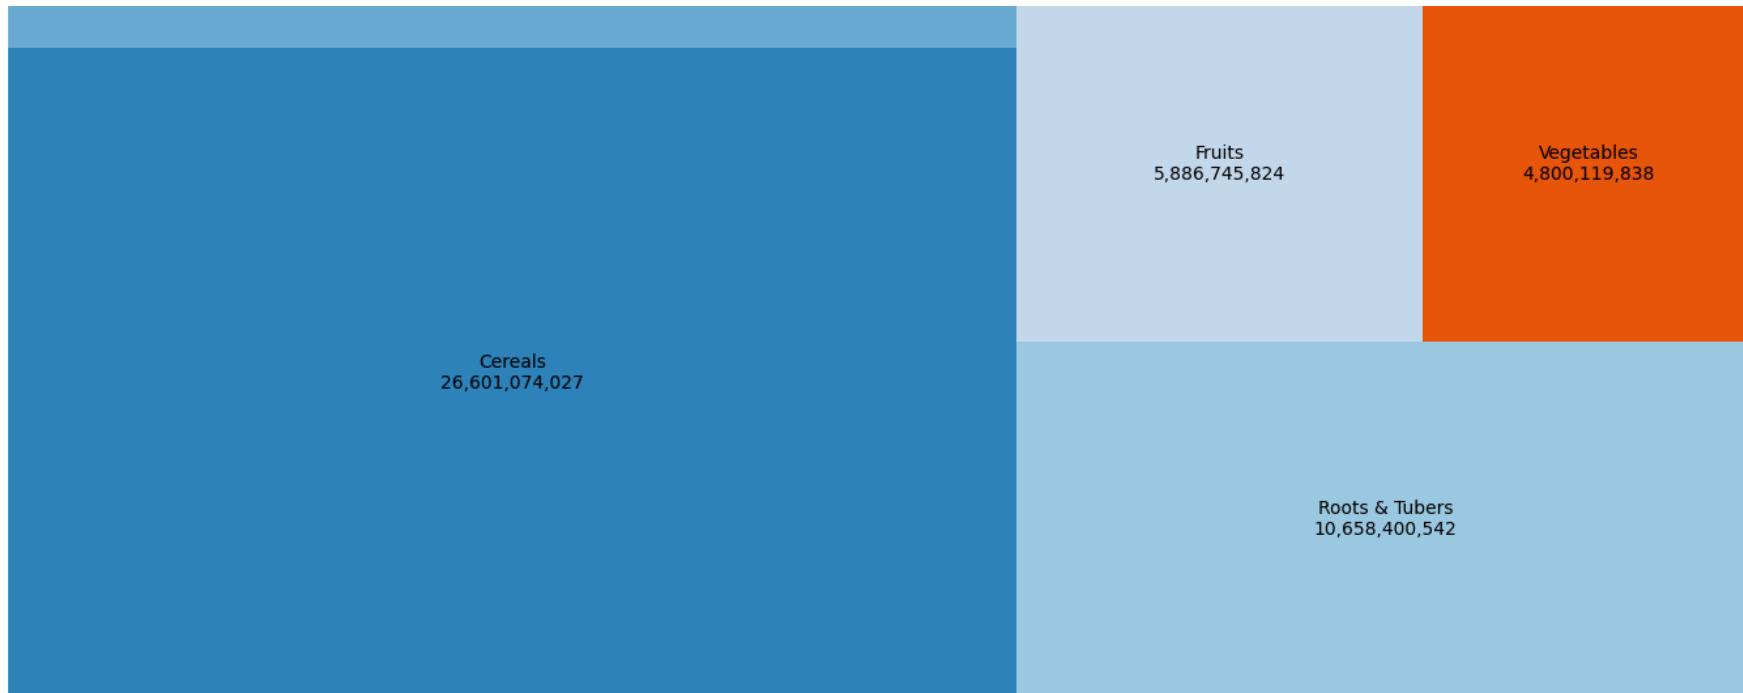
plt.show()
```

Treemap of Total Area Harvested by Group

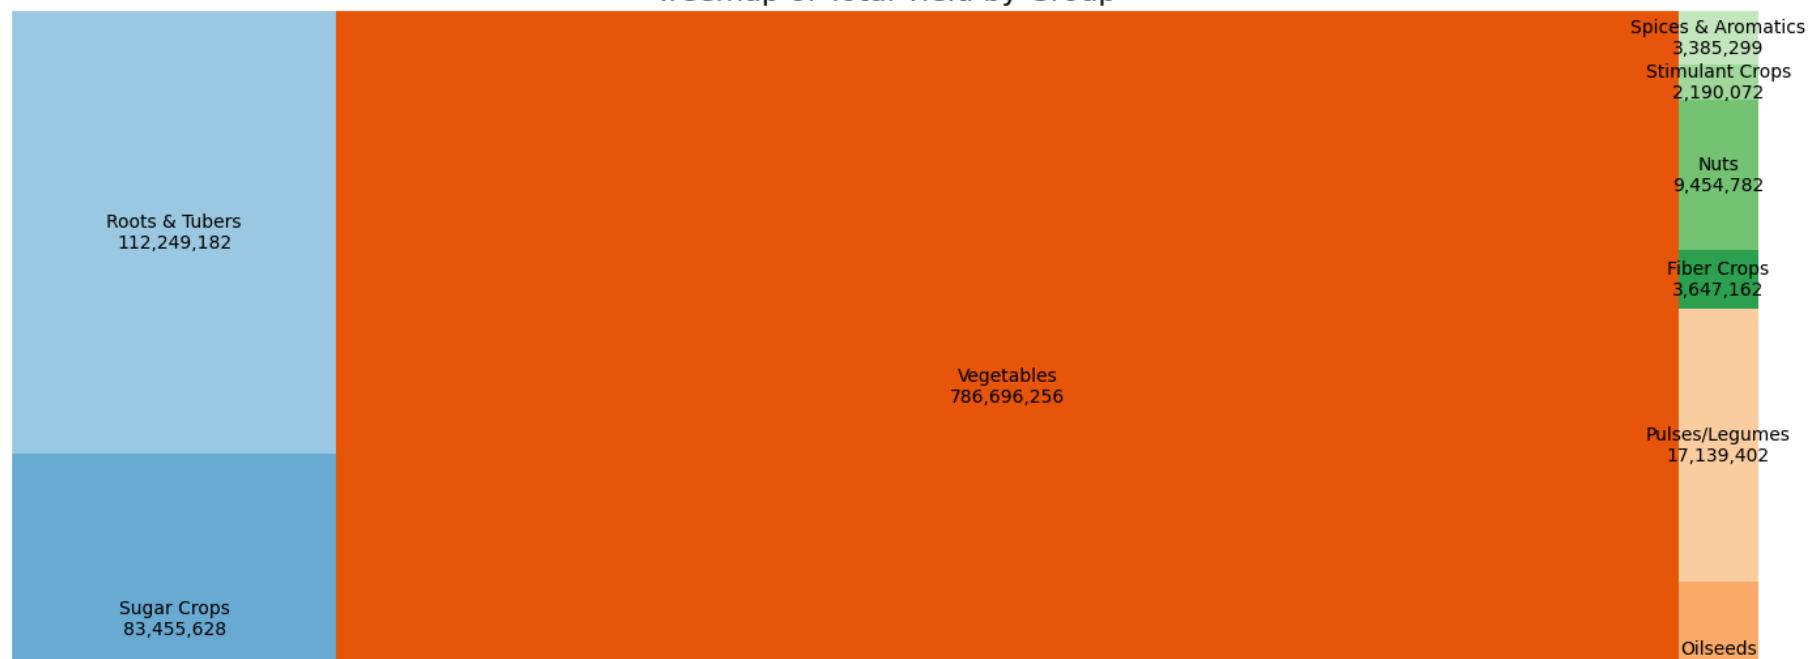


Treemap of Total Production by Group





Treemap of Total Yield by Group





```
In [96]: # 2) Define a helper function to map FAOSTAT names → ISO-3 codes
def faostat_to_iso3(name):
    try:
        return pycountry.countries.lookup(name).alpha_3
    except (LookupError, AttributeError):
        return None

# 3) Collect all three metrics into one “long” DataFrame for convenience
long_df = (
    df_merged
    .melt(
        id_vars=['Area', 'Year'],
        value_vars=['Production', 'Area_harvested', 'Yield'],
        var_name='Metric',
        value_name='Value'
    )
    # Drop any NaN values in Value (just in case)
    .dropna(subset=['Value'])
)

# 4) Lookup ISO3 for each row's “Area”
long_df['iso_alpha'] = long_df['Area'].apply(faostat_to_iso3)

# 5) Identify any unmapped areas (once) and override manually if needed
missing_iso = long_df.loc[long_df['iso_alpha'].isna(), 'Area'].unique()
if len(missing_iso) > 0:
    print("Unmapped FAOSTAT areas (will be dropped or overridden):", missing_iso)

# 6) Manual overrides for known FAOSTAT names that pycountry can't resolve:
manual_overrides = {
```

```

    "Belgium-Luxembourg":           None,   # drop
    "Czechoslovakia":              None,   # drop
    "Netherlands (Kingdom of the)": "NLD",   # Netherlands
    "Serbia and Montenegro":       None,   # drop
    "USSR":                         None,   # drop
    "Yugoslav SFR":                None   # drop
}

# 7) Apply overrides
long_df.loc[long_df['iso_alpha'].isna(), 'iso_alpha'] = (
    long_df.loc[long_df['iso_alpha'].isna(), 'Area']
    .map(manual_overrides)
)

# 8) Drop any rows still missing iso_alpha
long_df = long_df.dropna(subset=['iso_alpha'])

# 9) For each metric, build and display an animated choropleth
for metric_name, title, color_label in [
    ("Production",      "Crop Production by Country over Time (Europe, 1961-2023)",      "Production (t)"),
    ("Area_harvested",  "Crop Area Harvested by Country over Time (Europe, 1961-2023)",  "Area Harvested (ha)"),
    ("Yield",            "Crop Yield by Country over Time (Europe, 1961-2023)",            "Yield (kg/ha)")
]:
    # 9a) Filter the Long DF to just this metric
    df_metric = long_df[long_df['Metric'] == metric_name].copy()

    # 9b) Group & sum by (Area, Year, iso_alpha)–just in case there were multiple rows per country/year
    country_year = (
        df_metric
        .groupby(['Area', 'Year', 'iso_alpha'], as_index=False)[['Value']]
        .sum()
    )

    # 9c) Build the animated choropleth
    fig = px.choropleth(
        country_year,
        locations="iso_alpha",
        color="Value",
        hover_name="Area",
        animation_frame="Year",
        range_color=[country_year["Value"].min(), country_year["Value"].max()],

```

```
    color_continuous_scale="Viridis",
    title=title,
    scope="europe"
)

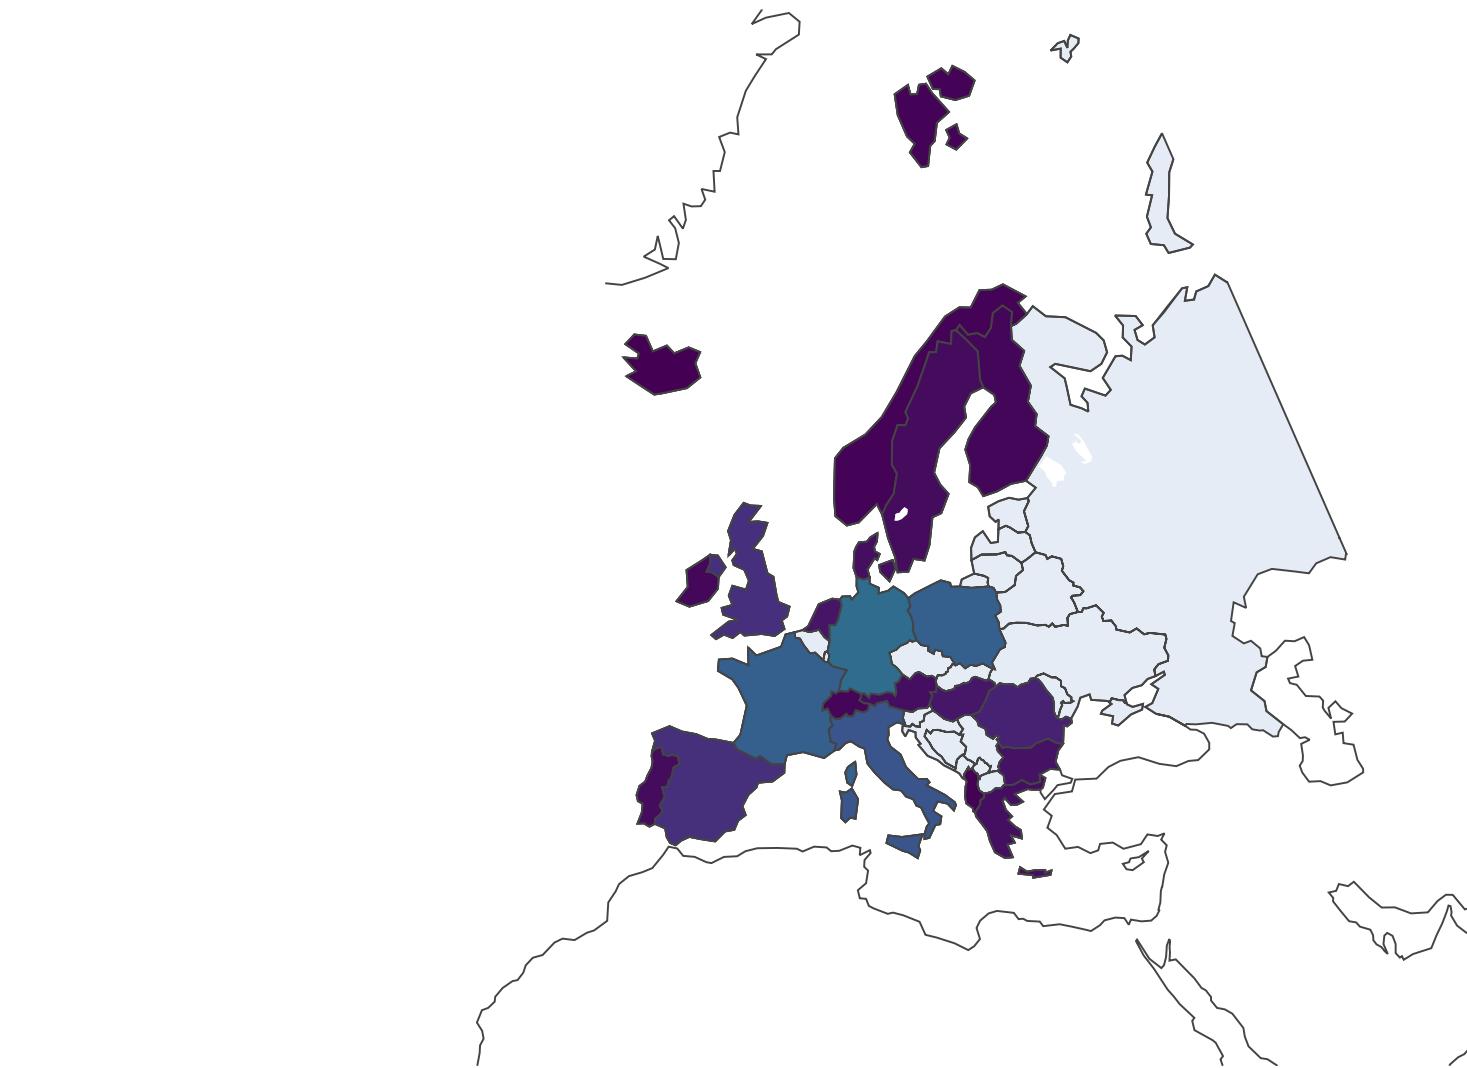
# 9d) Tweak Layout (size, remove frame)
fig.update_layout(
    width=1200,
    height=800,
    geo=dict(showframe=False, showcoastlines=True),
    title_x=0.5
)

# 9e) Change the colorbar Label
fig.update_coloraxes(colorbar_title=color_label)

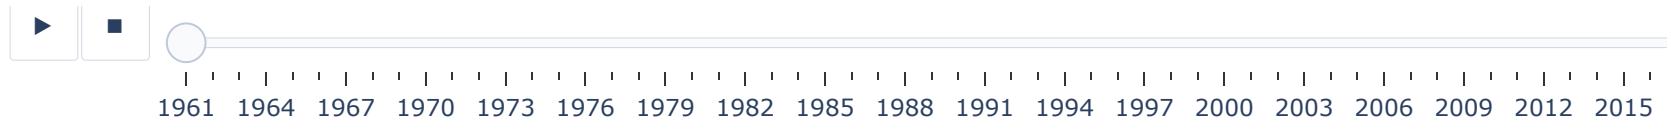
# 9f) Render as HTML in the notebook
html_str = fig.to_html(include_plotlyjs="cdn")
display(HTML(html_str))
```

Unmapped FAOSTAT areas (will be dropped or overridden): ['Belgium-Luxembourg' 'Czechoslovakia' 'Netherlands (Kingdom of the)' 'Serbia and Montenegro' 'USSR' 'Yugoslav SFR']

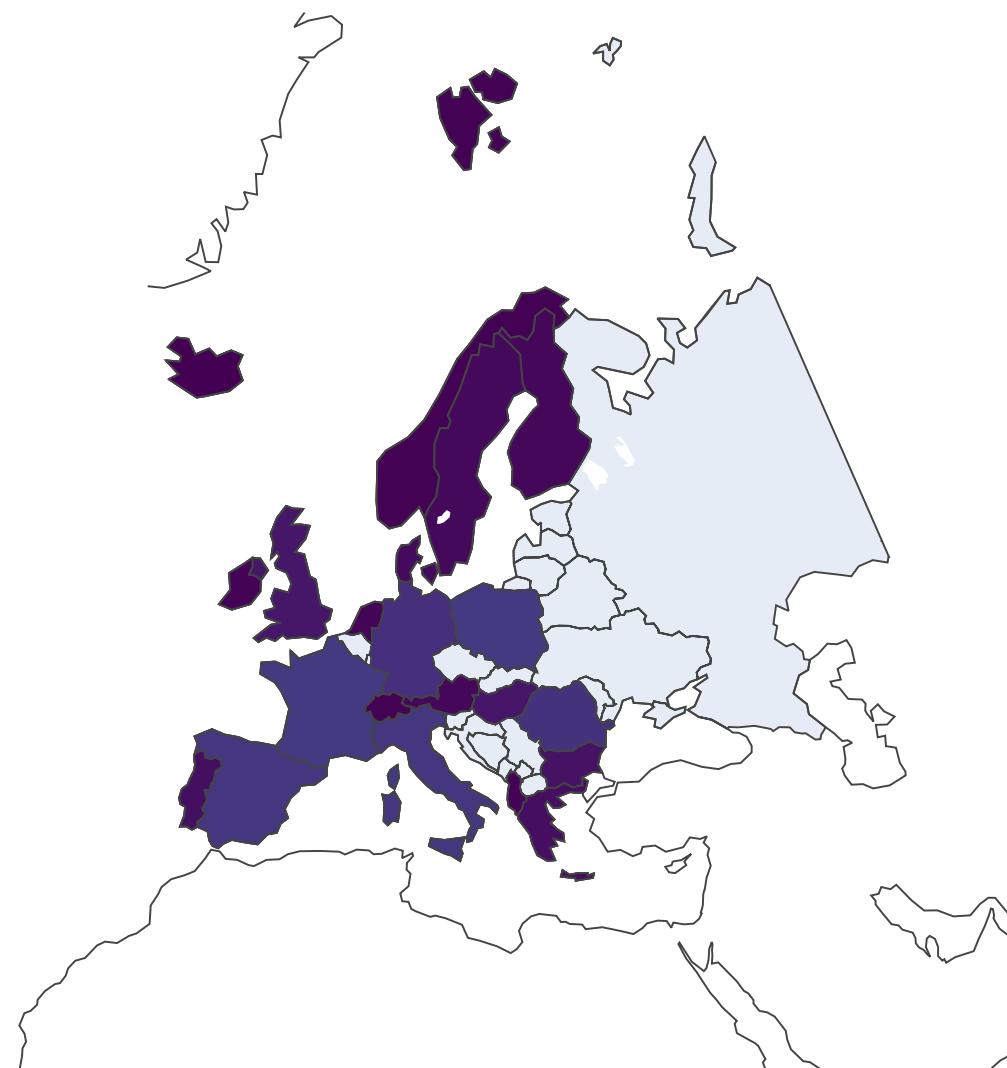
Crop Production by Country over Time (Europe, 1961–2023)



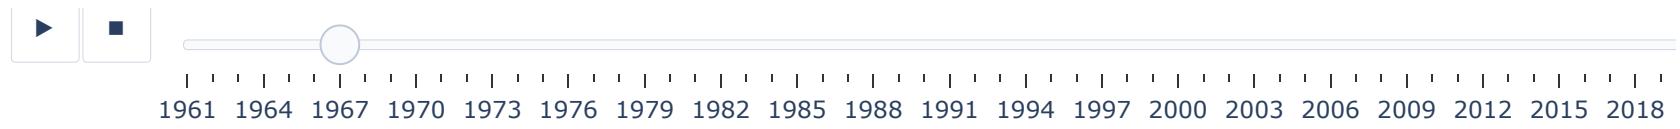
Year=1961



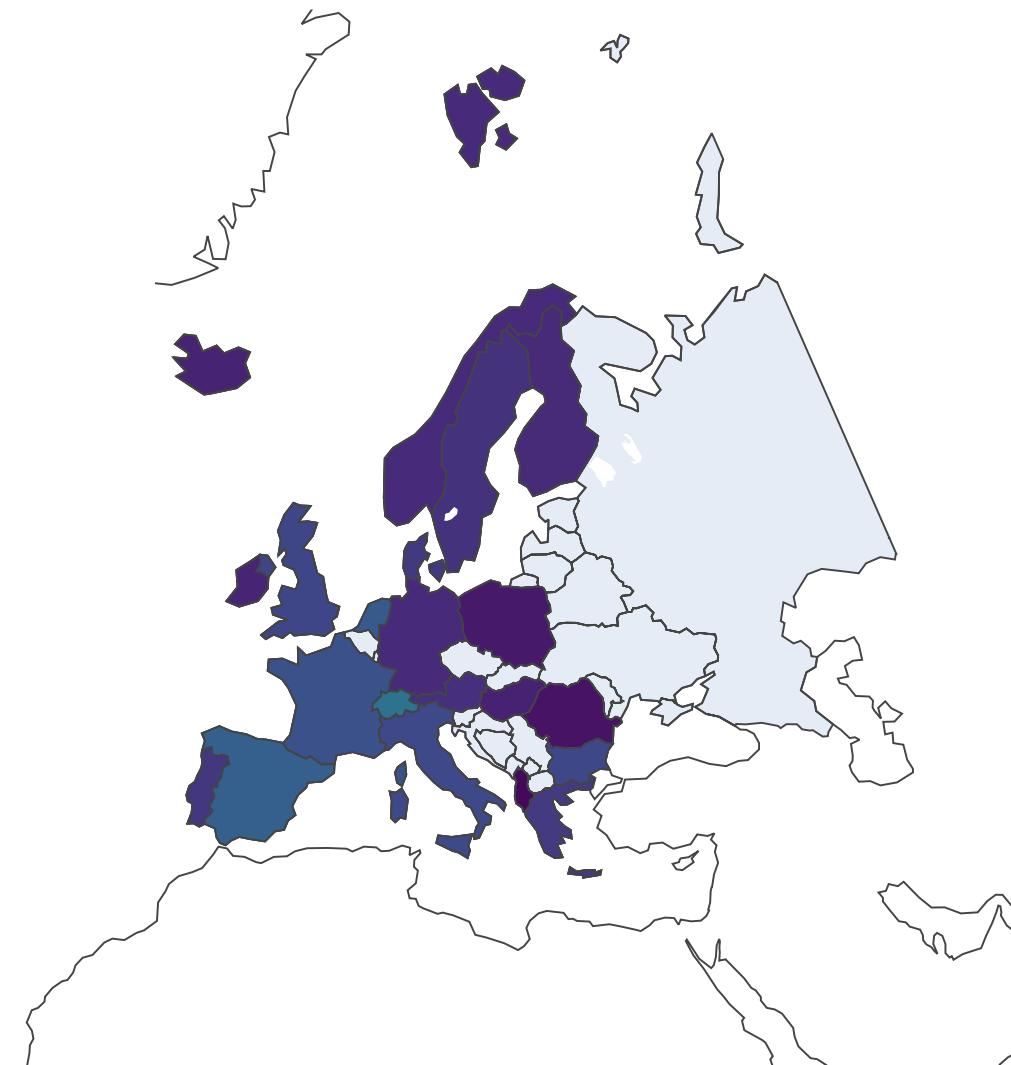
Crop Area Harvested by Country over Time (Europe, 1961–2023)



Year=1967



Crop Yield by Country over Time (Europe, 1961–2023)



Year=1967

