

BACKPRESSURE MECHANISMS IN STREAMING SYSTEMS KAFKA VS FLINK

Author: Turcu Ciprian-Stelian

THE BACKPRESSURE PROBLEM

What is Backpressure?

- Flow control mechanism when producers send data faster than consumers can process
- Without backpressure: system crashes, data loss, or memory overflow
- Critical for system stability under variable loads

Why It Matters in Big Data

- Billions of events per day in production systems
- Processing rates vary by operator complexity
- Temporary load spikes are inevitable
- Must maintain data integrity and low latency

Real-World Scenario

- Kafka topic receiving 1M events/sec
- Consumer can only process 500K/sec
- What happens to the other 500K events?

TWO DIFFERENT PHILOSOPHIES

1

Pull-Based Flow Control (Kafka)

- Consumer explicitly requests data at its own pace
- "I'll take records when I'm ready"
- No upstream feedback to producers
- Broker acts as durable buffer

2

Push-Based Backpressure (Flink)

- Framework automatically throttles data flow
- "Stop sending, my buffers are full"
- Backpressure propagates upstream automatically
- Built into network layer

KAFKA'S PULL-BASED APPROACH

How It Works

- 1.Consumer calls poll() to fetch records from broker
- 2.Consumer controls batch size with max.poll.records
- 3.Kafka topics store data durably on disk
- 4.Pull frequency determines consumption rate

Key Configuration Parameters

- max.poll.records: Records per poll
- max.poll.interval.ms: Max time between polls
- fetch.min.bytes: Minimum data before broker responds
- enable.auto.commit: Automatic vs manual offset commits

FLINK'S CREDIT-BASED BACKPRESSURE

How It Works

- 1.Credit-based flow control (similar to TCP sliding window)
- 2.Automatic propagation through entire pipeline
- 3.No application code needed

The Credit System

- Downstream announces: "I have N buffers available"
- Upstream sends backlog: "I want to send M messages"
- Downstream grants credits: "Send X buffers"
- Upstream sends only X buffers
- Each buffer sent = -1 credit


UBER AD EVENTS

System Requirements

- Process millions of ad events per second
- Exactly-once semantics end-to-end
- Sub-second latency for real-time decisions
- Multi-region deployment for resilience

Why Both Technologies

- Kafka: Durable buffering, handles producer spikes
- Flink: Complex stateful aggregations, exactly-once processing
- Kafka again: Decouples processing stages
- Final sinks: Pinot (real-time) + Hive (batch analytics)



NETFLIX RDG

Real-Time Distributed Graph

- Track member interactions across streaming, live events, games
- ~1 million messages/second per Kafka topic
- Near-real-time graph updates for personalization

Technology Choice

- Kafka: Event ingestion from Netflix apps
- Flink: "Strong capabilities around near-real-time event processing"
- Flink: Filtering, enrichment, graph transformations
- Each job independently tuned for its topic's characteristics

Architecture Evolution

- Initially: One Flink job consuming all Kafka topics
- Problem: "Became operational nightmare, tuning extremely difficult"
- Solution: 1-to-1 mapping of Kafka topic to Flink job
- Trade-off: More jobs but simpler operations

ADVANTAGES & TRADE-OFFS

Kafka Advantages

- Simple pull model, easy to understand
- Durable buffering survives failures
- Complete producer/consumer decoupling
- Language-agnostic consumers
- Mature operational tooling

Kafka Disadvantages

- Manual flow control implementation required
- Rebalancing storms from slow consumers
- max.poll.interval.ms false failures
- Consumer lag can grow unbounded
- Downstream systems can't signal back

Flink Advantages

- Automatic backpressure, no code needed
- True streaming, low latency
- Exactly-once guarantees
- Rich operators (windows, joins, aggregations)
- Built-in monitoring dashboard

Flink Disadvantages

- Operational complexity (cluster management)
- Memory-intensive for large state
- Steep learning curve
- Requires external durable source (Kafka)
- More components to manage

MONITORING IN PRODUCTION

Kafka Critical Metrics:

- consumer-lag: Offset difference
- records-consumed-rate: Throughput
- rebalance-count: Stability indicator
- fetch-latency-avg: Broker performance

Flink Critical Metrics:

- backPressuredTimeMsPerSecond: Time under backpressure
- idleTimeMsPerSecond: Waiting for data
- busyTimeMsPerSecond: Actual work time
- inPoolUsage/outPoolUsage: Buffer utilization


WHY USE BOTH TOGETHER

Kafka excels at:

- Durable message storage
- Decoupled communication
- Multi-consumer scenarios
- Inter-service messaging

Flink excels at:

- Complex stateful processing
- Low-latency transformations
- Windowing and temporal joins
- Exactly-once computations



Production Pattern:

- Kafka: Separation of concerns
- Flink: Complex transformations
- Kafka: Result distribution
- Both: Independent scaling

CONCLUSIONS & RECOMMENDATIONS

When to Choose Kafka Pull-Based

- Simple event-driven microservices
- Variable consumer processing rates
- Need maximum flexibility and decoupling
- Durability more important than millisecond latency

When to Choose Flink Credit-Based

- Complex stateful transformations required
- Sub-second latency critical
- Windowing, joins, aggregations needed
- Want automatic backpressure handling

When to Use Both

- Production Big Data platforms
- Need durability AND complex processing
- Kappa architecture (streaming-only)
- Scale requires separation of concerns

THANK YOU